

**Contact-Rich Planar Goal-Reaching:
Why Asymmetric Self-Play
Underperforms
Single-Agent Reward Shaping**
A Regime Analysis of a Method That Succeeds
Elsewhere

Vlad Iftime



**university of
 groningen**

**faculty of science
and engineering**



**university of
 groningen**

**faculty of science
 and engineering**

University of Groningen

**Contact-Rich Planar Goal-Reaching:
 Why Asymmetric Self-Play Underperforms Single-Agent Reward
 Shaping
 A Regime Analysis of a Method That Succeeds Elsewhere**

Master's Thesis

To fulfill the requirements for the degree of
 Master of Science in Artificial Intelligence
 at University of Groningen under the supervision of
 Prof. dr. L.R.B. Schomaker (Artificial Intelligence, University of Groningen)

Vlad Iftime (s3426394)

July 12, 2026

Contents

	Page
Abstract	6
1 Introduction	7
1.1 Contributions	8
1.2 Thesis Outline	9
2 Background and Related Work	10
2.1 Reinforcement Learning	10
2.1.1 Markov Decision Processes	10
2.1.2 Partial Observability and the POMDP	12
2.1.3 Model-Free and Model-Based Reinforcement Learning	12
2.2 Proximal Policy Optimization	13
2.3 Soft Actor-Critic	13
2.4 Potential-Based Reward Shaping	14
2.5 Curriculum Learning and Automatic Curricula	15
2.6 Asymmetric Self-Play	15
2.7 Goal-Conditioned and Hierarchical Reinforcement Learning	16
2.8 Mechanics and Geometry of Planar Pushing	17
2.9 Positioning of This Work	18
3 Methods	19
3.1 Simulation Environment	19
3.2 The Planar-Pushing Markov Decision Process	20
3.3 The Object-Relative Push Primitive	20
3.4 Why the Improvement Reward Underperforms	23
3.5 Potential-Based Reward Shaping for Pushing	23
3.6 The SE(2) Unified Distance	24
3.7 The Models	25
3.8 Why PBRS Assists Self-Play but Does Not Solve It	31
3.9 Experiment-Set Overview	31
3.10 Validation Protocol	33
3.11 Training Instrumentation for the Self-Play Mechanism	34
4 Experimental Setup	35
4.1 Tools and Infrastructure	35
4.2 Training Configuration and the Experiment Set	35
4.3 Validation Protocol	36
4.4 Cross-Environment Protocol	37
4.5 Off-Policy Baseline: SAC with Hindsight Experience Replay	37
4.6 Metrics	38

5	Results	39
5.1	RQ1: Environment-Efficiency of Reward Shaping	39
5.2	RQ2: Curriculum and Self-Play Against Single-Agent	40
5.3	RQ3: Skill Composition Under Self-Play	41
5.4	Cross-Environment Confirmation	42
5.5	Qualitative Behaviour and the Off-Policy Baseline	43
6	Discussion and Conclusion	45
6.1	Interpretation	45
6.2	Relation to Prior Work	45
6.3	Contributions Revisited	47
6.4	Limitations	47
6.5	Future Work	48
6.6	Conclusion	48
	Bibliography	49
	Appendices	53
A	Push Mechanics: Limit Surface and Characteristic Length	53
B	Training Curves	53
C	PBRs Parameter Sensitivity	54
D	Validation Scene Descriptions	54
E	cuRobo IK Pipeline and Workspace	54
F	Observation Space Specifications	54
G	Per-Test Error Bars	55
H	Curriculum Trigger Correction	55
I	Symmetric Self-Play Fail-Fast Equilibrium	56
J	Confounded Ablation Disaggregation	56
K	Training Budget Comparison	56
L	Secondary Variants	56

Abstract

On-policy reinforcement learning for contact-rich robotic manipulation is limited in practice by the amount of interaction it requires. This thesis studies a single goal-reaching task: planar pushing of a T-shaped block to a target pose by a simulated Universal Robots UR5e arm in NVIDIA Isaac Lab, under a fixed compute budget and one success criterion. The mechanics of pushing couple translation and rotation, so a goal that specifies both a target position and a target orientation cannot be decomposed into independent sub-problems. Asymmetric Self-Play (ASP) is an automatic curriculum that has succeeded on goal-reaching problems elsewhere, including robotic manipulation. On this multi-objective task, however, ASP underperforms a plain single-agent policy trained with a shaped reward, and this thesis identifies the regime conditions under which that occurs.

Nine models are compared under one protocol of 30 held-out scenes with best-of-20 stochastically-sampled trials. Three single-agent anchors establish that the task is solvable: a baseline with a hand-tuned improvement reward, a model with a Potential-Based Reward Shaping (PBRs) reward, and a forced position-to-orientation curriculum. Two self-play subjects, differing only in how the goal-proposing agent is rewarded, are the primary objects of study. Four further models isolate the cause: two repeat self-play on a rotationally symmetric disc whose goal constrains position alone, one uses a symmetric reward, and three cross-environment counterparts retrain the methods in a second simulator. PBRs combines bounded exponential potentials over position and orientation with a unified Special Euclidean group $SE(2)$ distance, and reaches single-agent success with roughly four times fewer parallel environments than the improvement reward. The final success rates are **TBD** until the seed runs are collated.

A per-axis analysis explains the self-play result. The self-play models learn the individual position and orientation skills, yet combine them far less often than two independent skills would predict. The coupled multi-objective gate, rather than the reward or the goal encoder, is the factor that separates self-play from the single-agent anchors. The disc contrast isolates this gate, and a replay diagnostic shows that the goal-proposing agent’s demonstrations are not reliably reproduced under stochastic contact. The result is scoped to this regime: ASP succeeds in the settings prior work used, and this thesis explains why it underperforms in contact-rich multi-objective goal-reaching under a single Graphical Processing Unit (GPU) budget. For this class of task, the practical priority is clear: investment in principled reward design produces a larger and more reliable improvement than investment in structured training dynamics.

1 Introduction

Service robots are expected to become part of everyday life, assisting people at home and at work with ordinary manipulation tasks. Meeting that expectation requires a robot to handle objects it has never seen, in cluttered and unstructured surroundings [1]. Bin picking is the canonical instance of this challenge: a robot must retrieve target items from a dense, disordered container, where objects rest against one another and few are graspable as they lie. Traditional pipelines that assume known object models and a clear grasp pose become brittle here, because the required information is often occluded or simply unavailable. Learning-based methods, and deep Reinforcement Learning (RL) in particular, have therefore become the dominant approach to manipulation in clutter [1].

A recurring insight in this literature is that grasping alone is not enough. When objects are packed together, the robot must first rearrange them, and non-prehensile actions such as pushing create the space and the poses that make a later grasp feasible. Push-grasp synergy, in which pushing and grasping are learned jointly, improves picking success in dense clutter beyond what grasping achieves on its own [2, 3, 4]. This positions pushing as a foundational pre-grasp skill for service robots: the ability to move an object into a target pose is what turns an ungraspable configuration into a graspable one. This thesis studies that skill in isolation, framed as a goal-reaching problem in which the agent must drive the object to a target pose. In doing so it works toward the same broad goal as the clutter-manipulation literature it builds on, namely competent, general manipulation in unstructured environments.

Progress on this goal is limited less by what RL can express than by what it costs to train. On-policy algorithms such as Proximal Policy Optimization (PPO) are stable and widely used [5], but they discard each batch of experience after a single update. Their sample requirement is therefore high, and it grows with the difficulty of the reward signal. When the reward is sparse or poorly shaped, most of the interaction budget yields little useful gradient, and the number of parallel environments needed for a stable update becomes the dominant practical constraint. Reducing this sample cost, and understanding which methods reduce it, is therefore a central concern for learning manipulation skills of this kind.

Two families of techniques aim to reduce this constraint, and they act on different parts of the problem. The first is reward shaping, which adds an auxiliary signal that directs the agent toward useful behaviour without changing what the task ultimately rewards. Potential-Based Reward Shaping (PBRS) is the principled form of this idea [6]: its shaping term is the difference of a state potential, which leaves the optimal policy unchanged. The second family modifies the training process rather than the reward. Curriculum learning orders tasks from easy to hard [7], and its automatic variants derive that order from the agent’s own competence. Asymmetric Self-Play (ASP) is the most prominent automatic form [8, 9]: one agent proposes goals while a second learns to reach them, so the goal distribution adapts to the learner without human specification. Both families are well motivated. Reward shaping, in the potential-based form used here, produces a single-agent policy that solves the task, and serves as the anchor of this study. ASP is the method under scrutiny: it has succeeded on goal-reaching problems elsewhere, including robotic manipulation [9] and reversible control tasks [8], yet on the task studied here it underperforms the single-agent anchor. That contrast is the puzzle this thesis addresses. A proven method underperforms a simpler one, and the aim is to explain why, by identifying the conditions of this task that differ from the settings in which the method succeeded.

The task studied here is planar pushing, the non-prehensile skill motivated above, reduced to a single object and framed as a goal-reaching problem so that the learning question can be

isolated. A robot must slide the object across a surface to a target pose. Pushing suits this study for a specific mechanical reason. Under the quasi-static assumption, a pushed object’s motion follows its friction limit surface, and the resulting twist couples translation and rotation [10, 11]. Almost every push therefore changes both the position and the orientation of the object. A goal that fixes both a target position and orientation is thus multi-objective, and the two objectives cannot be optimised independently. This coupling makes planar pushing a compact and demanding evaluation task: the reward must give an accurate gradient for two interacting quantities, and any curriculum or self-play mechanism must satisfy both criteria at once.

The experiments use a simulated Universal Robots UR5e arm with a Robotiq gripper, pushing a T-shaped block on a tabletop in NVIDIA Isaac Lab [12]. Isaac Lab runs many environment instances in parallel on a single Graphical Processing Unit (GPU), which lets every model train at a fixed compute budget on equal terms. The policy does not command joint torques directly; instead, it selects an object-relative push primitive with four parameters: an approach offset, an approach angle, a push length, and a push direction. A cuRobo Inverse Kinematics (IK) solver then executes the chosen push. Goals lie in the planar rigid-body configuration space $SE(2)$, and a push succeeds only when the object reaches a small position and orientation tolerance of the goal. Chapter 3 gives the full state, action, and reward definitions.

Against this background, the thesis poses three research questions, together with one supporting question on reward efficiency.

- **RQ1:** Does asymmetric self-play, as applied successfully in prior robotic manipulation, succeed in this contact-rich multi-objective goal-reaching task under a single-GPU budget?
- **RQ2:** Which difference between this task and the settings where self-play succeeded, among goal coupling, reward formulation, imitation signal, goal representation, and compute scale, accounts for the difference in outcome?
- **RQ3:** Where in the learning dynamics does self-play break down: the coupled success gate, a biased value function, a suppressed imitation signal, or demonstrations that cannot be reproduced under stochastic contact?
- **Supporting RQ:** Can potential-based reward shaping reach single-agent success with fewer parallel environments than a hand-tuned improvement reward?

The first question asks whether the method transfers to this regime at all, comparing self-play against the single-agent anchor under one budget and one success criterion. The second question decomposes the difference between this task and the successful settings into five candidate axes, and tests each in turn. The comparison between a rotationally symmetric disc and the T-block isolates the coupled goal as a single variable. The third question examines the mechanism, using per-axis competence, the value-function bias, the imitation signal, and a replay diagnostic that measures whether the goal-proposing agent’s own demonstrations reach the goals it sets. The supporting question isolates the contribution of reward shaping alone.

1.1 Contributions

This thesis makes three contributions, in order of priority. The first is the mechanism behind the self-play result. The models learn the individual position and orientation sub-skills, but do not

compose them under the coupled success gate; the coupled multi-objective gate, rather than the reward or the goal encoder, is the dominant cause. Five measurements support this: the disc-against-T-block contrast, a per-axis-against-combined diagnostic, a value-bias measurement, an imitation clip-saturation measurement, and a replay diagnostic. The second contribution is the regime contrast. It sets this task against the settings where self-play succeeded, along five axes: goal coupling, reward formulation, imitation signal, goal representation, and compute scale. A negative result here is thereby explained as a difference in regime, not a refutation of prior work. The third contribution is a reward-efficiency result. Potential-based reward shaping reaches single-agent success with roughly four times fewer parallel environments than the improvement reward, which anchors the task as solvable. The practical message follows: principled reward design should come before changes to the training dynamics. We scope every self-play result to this regime. Asymmetric self-play succeeds in the settings prior work used; the contribution here is the analysis of why it underperforms in this one.

1.2 Thesis Outline

The remainder of this thesis is organised as follows. Chapter 2 reviews the background and related work, covering RL and PPO, potential-based reward shaping, curriculum learning and self-play, goal-conditioned learning, and the mechanics of pushing. That chapter also draws the mechanical contrast between the disc and the T-block, and explains the multi-objective nature of the goal. Chapter 3 defines the task and its Markov Decision Process, explains why a naive improvement reward underperforms, and derives the potential-based reward and its SE(2) form. It then describes the nine models grouped by role, and sets out the experiment set that produces the results. Chapter 4 presents the experimental setup: the simulation and training configuration, the phases of the experiment set, the held-out validation protocol, the cross-environment comparison, and an off-policy baseline. Chapter 5 reports the results for each research question, namely whether self-play succeeds here, which regime axis accounts for the outcome, the underlying mechanism, and the environment-efficiency of shaping. Chapter 6 discusses the findings, states the limitations, and outlines future work. The appendices collect training curves, parameter sensitivity, scene descriptions, and detailed diagnostics.

2 Background and Related Work

This chapter develops the theory the thesis relies on and situates the work within the surrounding literature. It begins with the formal definitions of reinforcement learning, first for fully observable problems and then for the partially observable case that motivates the recurrent policy used here. It then describes the two algorithms used in the experiments, Proximal Policy Optimization and Soft Actor-Critic, before describing the two techniques the thesis compares: potential-based reward shaping, and structured training through curriculum learning and asymmetric self-play. The chapter ends with the mechanics and geometry of planar pushing, which explain why the task is multi-objective and why its two objectives cannot be optimised in isolation.

2.1 Reinforcement Learning

Reinforcement Learning (RL) is the branch of machine learning concerned with how an agent should act in an environment so as to maximise a cumulative scalar reward. Unlike supervised learning, the agent is never told the correct action; it receives only an evaluative signal and must discover, through trial and error, which behaviours lead to high long-term return. The agent observes the state of the environment, selects an action, and receives a reward together with a new state; repeating this loop produces a trajectory whose total reward the agent learns to improve. This interaction is formalised through the Markov Decision Process.

2.1.1 Markov Decision Processes

A Markov Decision Process (MDP) is the standard model of a fully observable sequential decision problem. Following [6, 13], it is a tuple $\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$, where:

- $\mathcal{S}$ is the **state space**, the set of all configurations the environment can occupy. At time step t the agent perceives a state $s_t \in \mathcal{S}$.
- $\mathcal{A}$ is the **action space**, the set of actions available to the agent; in robotics this may be a set of joint commands or, as in this thesis, a set of parameterised motion primitives.
- $P(s_{t+1} \mid s_t, a_t)$ is the **transition function**, giving the probability of reaching s_{t+1} after taking action a_t in state s_t . It encodes the environment dynamics.
- $R(s_t, a_t, s_{t+1})$ is the **reward function**, the scalar feedback for a transition.
- $\gamma \in [0, 1]$ is the **discount factor**, which sets the relative weight of immediate against future reward.

The defining property is the *Markov* assumption: the transition and reward depend only on the current state and action, not on the history that preceded them. Formally, $P(s_{t+1} \mid s_t, a_t) = P(s_{t+1} \mid s_0, a_0, \dots, s_t, a_t)$. The state is therefore a sufficient statistic of the past, which is what makes value functions and the Bellman recursion well defined.

The agent acts according to a **policy** $\pi(a \mid s)$, a mapping from states to a distribution over actions. Its objective is to maximise the **expected return**, the expected discounted sum of future rewards:

$$J(\pi) = \mathbb{E}_{\pi} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) \right]. \quad (1)$$

The discount factor γ has a dual role. Mathematically it keeps the infinite sum bounded; behaviourally it interpolates between a myopic agent ($\gamma \rightarrow 0$), which optimises only the immediate reward, and a far-sighted agent ($\gamma \rightarrow 1$), which weighs distant rewards almost as heavily as immediate ones. In episodic tasks such as the pushing problem studied here, the horizon is finite and the undiscounted case $\gamma = 1$ is admissible provided every trajectory eventually terminates, a subtlety that becomes important for reward shaping in Section 2.4.

Value functions. To reason about long-term consequences, RL introduces two value functions. The **state-value function** $V^\pi(s)$ is the expected return obtained by starting in state s and following π thereafter,

$$V^\pi(s) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) \mid s_0 = s \right], \quad (2)$$

while the **action-value function** $Q^\pi(s, a)$ conditions additionally on taking action a first,

$$Q^\pi(s, a) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) \mid s_0 = s, a_0 = a \right]. \quad (3)$$

The two are linked by $V^\pi(s) = \mathbb{E}_{a \sim \pi(\cdot|s)}[Q^\pi(s, a)]$, and their difference defines the **advantage** $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$, which measures how much better action a is than the policy's average behaviour in state s . The advantage is central to the policy-gradient methods used in this thesis.

Bellman equations. Value functions satisfy a recursive consistency condition. For the action-value function,

$$Q^\pi(s, a) = \mathbb{E}_{s' \sim P(\cdot|s, a)}[R(s, a, s') + \gamma V^\pi(s')], \quad (4)$$

which decomposes the value of a state-action pair into the immediate reward and the discounted value of the successor state. The optimal value functions $V^*(s) = \sup_\pi V^\pi(s)$ and $Q^*(s, a) = \sup_\pi Q^\pi(s, a)$ induce the optimal policy $\pi^*(s) = \arg \max_a Q^*(s, a)$. Whether solved by dynamic programming, value iteration, or sampled approximation, essentially all RL algorithms can be seen as ways of estimating or improving the quantities in Equation 4.

Exploration versus exploitation. Because the agent learns from its own experience, it faces a fundamental trade-off. *Exploitation* chooses the action currently believed to be best, which secures known reward; *exploration* tries alternatives to improve those beliefs, at the risk of short-term loss for long-term gain. An agent that only exploits may converge to a suboptimal behaviour, while one that only explores never applies what it has learned. Mechanisms such as ϵ -greedy action selection and entropy regularisation, the latter central to Soft Actor-Critic (Section 2.3), exist to manage this trade-off. In sparse-reward manipulation the trade-off is difficult, because informative reward is encountered so rarely that undirected exploration is unlikely to reach it, a difficulty that both reward shaping and automatic curricula attempt to address.

2.1.2 Partial Observability and the POMDP

The MDP assumes the agent perceives the full state. Real and simulated robots rarely enjoy this: sensors are noisy, occlusions hide part of the scene, and quantities such as contact forces or friction coefficients are simply not observable. Such problems are modelled as a Partially Observable Markov Decision Process (POMDP), a tuple $(\mathcal{S}, \mathcal{A}, P, R, \Omega, O, \gamma)$ that augments the MDP with an **observation space** Ω and an **observation function** $O(o | s', a)$ giving the probability of observation o after a transition to s' . The agent no longer sees s_t ; it sees only o_t , and a single observation is generally not a sufficient statistic of the state.

Because the observation breaks the Markov property, an optimal policy cannot in general be a function of the current observation alone. The classical solution is to act on a **belief state**, the posterior $b_t(s) = P(s_t = s | o_0, a_0, \dots, o_t)$ over states given the entire interaction history, which is itself Markovian; maintaining the exact belief, however, is intractable for continuous state spaces. A practical alternative, and the one adopted in this thesis, is to have the policy summarise the history with a recurrent network. A Long Short-Term Memory (LSTM) network maintains a hidden state that is updated at each step, which provides the policy with a learned, compressed approximation of the belief. In the pushing task the current observation reports object and goal poses, but not the underlying contact state or the outcome of prior pushes. The problem is therefore formally a POMDP. Conditioning the policy on the sequence of pushes through an LSTM lets it combine information across the episode, rather than respond to a single frame. For this reason every policy compared in this thesis is recurrent.

2.1.3 Model-Free and Model-Based Reinforcement Learning

RL algorithms divide into two families according to whether they build an explicit model of the environment dynamics $P(s_{t+1} | s_t, a_t)$ [14]. **Model-free** methods learn a policy or value function directly from sampled interaction, without ever estimating the transition function. They subdivide further:

- **Value-based** methods, such as Deep Q-Networks [15], learn an action-value function and act greedily with respect to it, obeying the recursion $Q(s, a) = \mathbb{E}[R(s, a, s') + \gamma \max_{a'} Q(s', a')]$.
- **Policy-based** methods, such as PPO [5], parameterise the policy directly and ascend the gradient of the expected return in Equation 1.
- **Actor-critic** methods, such as Soft Actor-Critic [16], combine the two, maintaining an actor that represents the policy and a critic that estimates value to reduce the variance of the policy gradient.

Model-free methods make no assumptions about the dynamics and are robust to complex, contact-rich physics, which makes them the default choice in manipulation; their weakness is high sample complexity, since every unit of information must be extracted from interaction.

Model-based methods instead learn an approximate dynamics model $\hat{P}(s_{t+1} | s_t, a_t)$ and use it to plan or to generate synthetic experience, typically by minimising a prediction error such as

$$L_{\text{model}} = \mathbb{E}_{\mathcal{D}} \left[\left\| P(s_{t+1} | s_t, a_t) - \hat{P}(s_{t+1} | s_t, a_t) \right\|^2 \right], \quad (5)$$

over a dataset $\mathcal{D}$ of observed transitions. They can be far more sample-efficient, but their performance is bounded by the fidelity of the learned model, and small errors compound over

long rollouts. The choice between the two families is therefore a choice between robustness and sample efficiency. This thesis works entirely in the model-free, on-policy regime and asks a complementary question: rather than learning a dynamics model to save samples, can a principled reward do so instead, by making each sampled transition more informative? Section 2.4 takes up this question.

2.2 Proximal Policy Optimization

Proximal Policy Optimization (PPO) [5] is the on-policy, model-free algorithm used to train every policy in this thesis. It belongs to the policy-gradient family, which optimises a parameterised policy $\pi_\theta(a \mid s)$ by ascending an estimate of the gradient of the expected return. The basic policy-gradient estimator,

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\nabla_\theta \log \pi_\theta(a_t \mid s_t) \hat{A}_t \right], \quad (6)$$

scales the score function of each action by its estimated advantage $\hat{A}_t$, so that actions better than average are made more likely. In practice the advantage is computed with Generalized Advantage Estimation (GAE), which trades bias against variance by exponentially weighting multi-step temporal-difference errors; GAE reappears in Chapter 3, where its decomposition explains how reward shaping interacts with a stale value estimate.

PPO improves on the earlier Trust Region Policy Optimization [17], which constrained each update to a trust region through a costly second-order optimisation. PPO achieves a similar effect with a first-order **clipped surrogate objective**. Writing the probability ratio between the new and old policies as $r_t(\theta) = \pi_\theta(a_t \mid s_t) / \pi_{\theta_{\text{old}}}(a_t \mid s_t)$, the objective is

$$L^{\text{CLIP}}(\theta) = \mathbb{E}_t \left[\min \left(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]. \quad (7)$$

The clipping removes the incentive to move the policy ratio beyond $[1 - \epsilon, 1 + \epsilon]$, which keeps each update inside an implicit trust region and prevents the large, destabilising steps that occur in unconstrained policy gradients. This same clipping mechanism is later reused in the behavioural-cloning loss of asymmetric self-play (Section 2.6), where it has an unintended consequence that this thesis diagnoses. The full PPO objective adds a value-function loss $L^{\text{VF}}(\phi) = \mathbb{E}_t[(V_\phi(s_t) - \hat{R}_t)^2]$ and an entropy bonus $L^{\text{ENT}} = \mathbb{E}_t[\mathcal{H}(\pi_\theta(\cdot \mid s_t))]$ that encourages exploration:

$$L(\theta) = L^{\text{CLIP}}(\theta) - c_1 L^{\text{VF}}(\phi) + c_2 L^{\text{ENT}}, \quad (8)$$

with coefficients c_1 and c_2 balancing the three terms.

2.3 Soft Actor-Critic

Soft Actor-Critic (SAC) [16] is an off-policy actor-critic algorithm built on the maximum-entropy framework. Where standard RL maximises return alone, SAC augments the objective with the entropy of the policy,

$$J(\pi) = \sum_{t=0}^T \mathbb{E}_{(s_t, a_t) \sim \rho_\pi} \left[R(s_t, a_t) + \alpha \mathcal{H}(\pi(\cdot \mid s_t)) \right], \quad (9)$$

where $\mathcal{H}(\pi(\cdot \mid s_t)) = \mathbb{E}[-\log \pi(\cdot \mid s_t)]$ is the policy entropy and the temperature α sets the weight of exploration against reward. Maximising entropy encourages the policy to remain as

stochastic as the reward permits. This improves exploration and robustness, and discourages premature convergence to a single mode, an advantage in sparse-reward or multimodal problems. Being off-policy, SAC reuses past experience from a replay buffer, and is typically far more sample-efficient than on-policy methods. SAC is relevant here as the basis of an off-policy comparison baseline. Combined with Hindsight Experience Replay [18], which relabels failed trajectories with the goals they happened to achieve, it provides a natural goal-conditioned learner against which the on-policy results can be calibrated, as described in Chapter 4.

2.4 Potential-Based Reward Shaping

Reward shaping supplies an agent with additional reward beyond the sparse task signal, to guide learning through the temporal credit-assignment problem [19]. Its danger is equally well known: an ill-chosen shaping reward can change the optimal policy, so that the agent learns to exploit the shaping rather than to solve the task. [6] give the canonical illustration, a bicycle agent rewarded for progress toward a goal that learns to ride in circles near the start, collecting the progress reward repeatedly without ever advancing.

Potential-Based Reward Shaping (PBRs) is the principled resolution of this danger. [6] prove that a shaping reward of the form

$$F(s, a, s') = \gamma \Phi(s') - \Phi(s), \quad (10)$$

defined as the discounted difference of a **potential function** $\Phi : \mathcal{S} \rightarrow \mathbb{R}$ over states, leaves the optimal policy of the MDP unchanged for *any* choice of Φ . The intuition is that the shaping term telescopes along a trajectory: summed over an episode it collapses to a boundary term that depends only on the first and last states, so it cannot create the reward-collecting cycles that non-potential shaping permits. Formally, adding Equation 10 to the task reward shifts every action-value by the same state-dependent constant, $Q'(s, a) = Q(s, a) - \Phi(s)$, and since the shift is independent of the action it does not alter the arg-max that defines the optimal policy. This decoupling is the property the thesis exploits: it permits a dense, informative reward to be designed purely to accelerate learning, with a guarantee that the task being solved is unchanged.

Subsequent work extended the framework in directions this thesis uses. [20] generalised the potential to depend on time as well as state, $F(s, t, s', t') = \gamma \Phi(s', t') - \Phi(s, t)$. This preserves policy invariance while letting the shaping adapt during learning, and so licences shaping that changes across a curriculum. [21] showed that an arbitrary reward function can be expressed as potential-based advice, widening the class of admissible potentials. Crucial for the present setting is the episodic case. [22] and [19] analysed PBRs in the finite-horizon setting and clarified when the undiscounted case $\gamma = 1$ is safe. Provided every trajectory terminates and terminal states are handled consistently, the shaping sum telescopes to $\Phi(s_T) - \Phi(s_0)$ and remains bounded. This result justifies the use of $\gamma_{\text{shaping}} = 1$ in the potentials constructed in Chapter 3. The theorem and its episodic refinement are the foundation of the reward design in this thesis. The specific potentials, which combine bounded exponential functions of position and orientation error, are introduced there rather than here, since they are a contribution of the work rather than prior art.

2.5 Curriculum Learning and Automatic Curricula

An orthogonal approach to sample inefficiency is to change not the reward but the order in which the agent encounters tasks. Curriculum learning organises training from easy to hard, so that competence acquired on simple problems supports progress on difficult ones [7]. The survey of [7] provides a taxonomy that distinguishes, among other axes, *forced* curricula, in which a designer fixes the sequence of tasks or objectives in advance, from *automatic* curricula, in which the sequence is derived from the learner’s own progress. Both appear in this thesis: a forced position-then-rotation curriculum, and an automatic curriculum generated by self-play. Several forced and semi-automatic schemes inform the design. Reverse curriculum generation [23] starts the agent near the goal and expands the start distribution outward as competence increases, so that the agent always trains at the boundary of what it can already achieve. Precision-based curricula such as Continuous Curriculum Learning [24] progressively tighten the success tolerance, which is a close analogue of staging a pushing task from coarse positioning to fine orientation control. The survey of [25] organises automatic curriculum methods into families that include self-paced learning, teacher-student setups, and goal-generation, and notes the recurring difficulty that an automatic curriculum can propose goals faster than the learner masters them. That failure mode is the one the self-play experiments in this thesis encounter.

2.6 Asymmetric Self-Play

Asymmetric Self-Play (ASP) is the automatic curriculum at the centre of the thesis. Introduced by [8], it pits two policies embodied in the same agent against each other: Alice proposes a task by acting in the environment, and Bob must then solve it. Because Alice demonstrates each task by performing it, every proposed goal is achievable by construction, and the competition between the two produces a curriculum of increasing difficulty with no external reward or human specification. Sukhbaatar et al. define a self-regulating reward structure,

$$R_B = -\gamma t_B, \quad (11)$$

$$R_A = \gamma \max(0, t_B - t_A), \quad (12)$$

where t_A is the time Alice takes to set the task, t_B the time Bob takes to solve it, and γ a scaling coefficient. Alice is rewarded when Bob is slow but penalised for her own effort, so her optimal behaviour is to find the *simplest* tasks that Bob cannot yet solve, placing each new task just beyond his current ability. This time-based reward is one of the two Alice reward formulations the thesis evaluates.

[9] scaled ASP to robotic manipulation and introduced the second formulation. They model the problem as a goal-augmented MDP $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \mathcal{G} \rangle$ with an explicit goal space $\mathcal{G}$, run Alice to produce a trajectory whose final state s_T becomes the goal for Bob, and train Bob from the same initial state on the resulting goal distribution $\mathcal{G}(g \mid s_0)$. Their Alice reward is *outcome-based* rather than time-based: Alice receives a reward when Bob fails to reach the goal, a shaping bonus for proposing a valid goal, and a penalty for placing objects out of bounds, while Bob receives a sparse reward for reaching the goal. Both the outcome-based and the time-based Alice rewards are compared in this thesis.

A distinctive feature of the manipulation formulation is that Alice’s own trajectory is itself a demonstration of the goal she sets, which [9] use through Alice Behavioral Cloning (ABC). Bob is trained on a combined objective $\mathcal{L} = \mathcal{L}_{\text{RL}} + \beta \mathcal{L}_{\text{abc}}$, where the behavioural-cloning term imitates Alice’s relabelled trajectory on goals Bob failed to solve. To prevent the imitation

term from producing destabilising policy updates, they apply the PPO clipping mechanism of Equation 7, with the advantage fixed to one:

$$\mathcal{L}_{\text{abc}} = -\mathbb{E}_{(s_t, g_t, a_t) \in \mathcal{D}_{\text{BC}}} \left[\text{clip} \left(\frac{\pi_B(a_t \mid s_t, g_t; \theta)}{\pi_B(a_t \mid s_t, g_t; \theta_{\text{old}})}, 1 - \epsilon, 1 + \epsilon \right) \right]. \quad (13)$$

This clipped behavioural-cloning loss behaves as intended only when Alice’s demonstrated actions lie close to Bob’s current policy. When the two diverge strongly, as they do under the discrete, contact-rich action space of this thesis, the clip saturates and the imitation gradient vanishes; this interaction is analysed in the results. Behavioural cloning more broadly provides the wider context for ABC. This includes cloning from observation [26], implicit energy-based cloning [27], and demonstration-boosted RL [28, 29], all methods for adding a demonstration signal to an RL objective.

The broader set of automatic curricula based on adversarial goal generation places ASP among several related methods. PAIRED [30] generates environments that maximise the regret between a protagonist and an antagonist. AMIGo [31] trains a teacher to propose goals a student finds challenging but achievable. Adversarial Intrinsic Motivation [32] directs exploration with a Wasserstein distance to the goal distribution. All share ASP’s central premise, that a well-tuned adversary produces a useful curriculum, and all share its central risk, that the adversary may generate tasks the learner cannot use. The stability of such coupled optimisation is itself a subject of study. [33] decompose the dynamics of differentiable games into potential and Hamiltonian components, and identify when adversarial training converges or cycles. [34] demonstrate that competitive self-play can reach very high performance, but only at very large scale. These two observations define the scope of the present study: the self-play results here are obtained under a single-Graphical-Processing-Unit budget, and are interpreted accordingly. The goal criterion also separates this thesis from the settings in which ASP succeeded, and the difference determines what the method must achieve. In the original formulation of [8], a goal is a single state Bob must reach, often in a domain where position dominates and orientation is absent or loosely constrained. In the manipulation setting of [9], a goal is the final pose of Alice’s trajectory, scored within a tolerance on each component. The task studied here instead applies a strict combined gate. A goal counts as reached only when position and orientation both lie within tolerance at the same time. The contact mechanics of Section 2.8 couple the two, so neither can be optimised alone. The goal is therefore genuinely multi-objective in a sense the prior settings did not require. This is a difference in regime, not a defect in the earlier work, and it defines one of the axes along which this thesis explains why a proven method underperforms here.

2.7 Goal-Conditioned and Hierarchical Reinforcement Learning

ASP produces a goal-conditioned policy, and the way a goal is represented influences what the policy can learn. Goal-conditioned RL augments the state with a goal and learns a single policy $\pi(a \mid s, g)$ that generalises across goals; Hindsight Experience Replay [18] makes this practical under sparse reward by relabelling achieved states as goals. Sukhbaatar et al. in their goal-embedding work [35] showed that compressing goals into a learned latent through a goal encoder can improve hierarchical learning, an idea the thesis adopts in the form of the GoalEncoder used by Bob. The effect of such a compression can be positive or negative, and the information-bottleneck literature explains why. [36] formalise the trade-off between compression and predictive sufficiency. InfoBot [37] applies it in RL, and shows that a bottleneck can either

improve generalisation by discarding distractors or reduce precision by discarding task-relevant detail. Whether the goal encoder improves or reduces performance in the pushing task is one of the questions the ablations address.

Hierarchical RL provides further context for the two-level Alice-Bob structure. Methods such as HIRO [38] and FeUdal Networks [39] decompose control into a high-level policy that sets sub-goals and a low-level policy that achieves them. The Option-Critic architecture [40] learns temporally extended actions with their own termination conditions, a useful lens on the push primitive used here as a macro-action. The survey of [41] situates these approaches and their open challenges. Imagined-goal methods [42] generate goals from a learned latent, mirroring how Alice’s proposals define Bob’s training distribution.

2.8 Mechanics and Geometry of Planar Pushing

The final body of theory explains why the pushing task is genuinely multi-objective. When a robot slides an object across a surface, the motion is governed by friction between the object and the support. The analysis is made tractable by the **quasi-static** assumption: pushing is slow enough that inertial forces are negligible and the applied load is balanced entirely by friction [10]. Under this assumption the object moves only while it is being pushed and stops the instant the push stops, so its trajectory is determined by the geometry of frictional contact rather than by dynamics.

The central tool is the **limit surface** [11]. At a sliding contact, the set of all friction forces and moments the interface can sustain forms a convex surface in the three-dimensional load space of planar forces and moment, $P = [F_x, F_y, M]$. For isotropic Coulomb friction this surface is a convex, closed shape often approximated by an ellipsoid. Its importance lies in the **normality (maximum-power) principle**: the object’s instantaneous motion twist $\mathbf{t} = [v_x, v_y, \omega]$, comprising a linear velocity and an angular velocity, is normal to the limit surface at the point corresponding to the applied load,

$$\mathbf{t} \propto \nabla f(\mathbf{w}), \quad (14)$$

where $f(\mathbf{w}) = 0$ is the limit surface and $\mathbf{w}$ the applied wrench. The consequence is important for this thesis. Consider any push not directed exactly through the object’s centre of friction. Its wrench lies where the surface normal has a non-zero moment component, so the resulting twist contains an angular velocity: *almost every push rotates the object as well as translating it*. Translation and rotation are mechanically coupled, and a goal that constrains both cannot be decomposed into two independent sub-problems. This coupling is the physical origin of the combined position-and-orientation success criterion that recurs throughout the results. It is also the reason a reward or curriculum that treats the two objectives as separable is inconsistent with the underlying physics. [43] developed the controllability and planning consequences of this mechanics for stable pushing, and [44] gave practical force-motion models. The survey of [45] and the high-fidelity dataset of [46] document how strongly real pushing behaviour depends on object shape and pressure distribution.

This dependence on object shape distinguishes the two objects used in this thesis and motivates one of its experiments. A rotationally symmetric disc has a radially symmetric pressure distribution about its centre. A push whose line of action passes through that centre applies a wrench with no moment component, so the limit-surface normal points opposite the push direction and the resulting twist is a pure translation. The disc therefore has no orientation to control, and a goal on the disc constrains position alone. The T-block has an asymmetric mass

distribution, and few pushes pass through its centre of friction; each off-centre push carries a moment, and by the normality condition of Equation 14 the resulting twist contains an angular velocity. Orientation is then an independent quantity that must be controlled alongside position. This mechanical difference is the basis of the coupling-isolation experiment of Chapter 3. The disc removes the orientation objective at the level of the physics. A self-play model that reaches the single-agent level on the disc, yet underperforms it on the T-block, thereby identifies the coupled objective, rather than the object itself, as the source of the difficulty.

Because position and orientation are coupled, it is natural to measure goal error with a single metric on the space of planar poses, rather than with two separate quantities. The configuration of a planar rigid body is an element of the **Special Euclidean group** $SE(2)$, the group of rigid-body transformations of the plane. Its elements combine a translation $(x, y) \in \mathbb{R}^2$ with a rotation $\theta \in SO(2)$. [47] formulated rigid-body kinematics and dynamics in the language of Lie groups, which provides the tools to define a distance on this space. A left-invariant Riemannian metric on $SE(2)$ yields a geodesic distance between two poses. Because translation is measured in metres and rotation in radians, combining them requires a **characteristic length** L that converts an angular error into an equivalent linear displacement. A pose error can then be summarised by a single scalar of the form

$$d_{\text{pose}} = \sqrt{\Delta x^2 + \Delta y^2 + L^2 \Delta \theta^2}, \quad (15)$$

in which L sets the relative weight of orientation against position and is naturally identified with a physical length scale of the object. This unified $SE(2)$ distance underlies the single-potential reward used for the self-play variants in Chapter 3. The idea of respecting the geometry and symmetry of the configuration space also appears in recent manipulation learning: $SE(3)$ -equivariant and diffusion-based methods [48], MDP homomorphic networks [49], equivariant transporter networks [50], and equivariant RL under partial observability [51] all exploit the group structure of the task, and provide broader context for a geometry-aware treatment of planar pushing.

2.9 Positioning of This Work

The literature provides two well-founded but rarely compared techniques for reducing the sample cost of on-policy manipulation learning: shaping the reward, with the policy-invariance guarantee of PBRS, and structuring the training process, through forced or automatic curricula such as ASP. Each has strong theoretical support and prominent empirical results, yet the two are seldom evaluated directly against each other on a single task, under a common budget, and against a common success criterion. The planar-pushing setting makes the comparison more stringent, because the limit-surface coupling of Equation 14 makes the task irreducibly multi-objective and therefore places a strong demand on any method that implicitly assumes the two objectives can be handled separately. The chapters that follow implement both techniques in this setting and measure what each contributes.

3 Methods

This chapter defines the goal-reaching task and the nine models that are compared on it. It begins with the simulation environment and the formal Markov Decision Process, including the object-relative push primitive that serves as the action. It then explains why the hand-tuned improvement reward of an initial baseline underperforms, and derives the potential-based reward that replaces it. The reward is presented first as separate position and orientation potentials, then as a single Special Euclidean group $SE(2)$ distance. The chapter closes with the models, grouped by the role each one serves, and with an analysis of why potential-based shaping partially assists self-play without matching single-agent training.

3.1 Simulation Environment

All experiments are performed in NVIDIA Isaac Lab 2.3.0, built on Isaac Sim 5.1.0, using the PhysX physics engine on the Graphical Processing Unit (GPU) [12, 52]. The agent is a Universal Robots UR5e arm fitted with a Robotiq 2F-140 parallel-jaw gripper. The manipulated object is a T-shaped block resting on a tabletop, and the goal is displayed as a coloured ghost of the same block at the target pose. Figure 1 shows the setup.




Figure 1: The planar-pushing task. A UR5e arm pushes a T-block (blue) across a tabletop toward a goal pose (orange). Goals are defined in $SE(2)$ and combine a target position and a target orientation.

The principal advantage of Isaac Lab for this study is GPU-batched parallelism. Many copies of the environment are stepped at the same time on a single device, which produces the large sample batches on-policy learning requires. The single-agent and curriculum models are trained with 528 parallel environments, which yields a batch of roughly 7,920 push transitions per Proximal Policy Optimization (PPO) update. A hardware upper bound of one RTX Pro 6000 GPU (96 GB of memory) sets the compute budget for every model, so all methods are compared on equal terms. Inverse Kinematics (IK) for the arm is solved with cuRobo 0.7.5, a GPU-accelerated batched IK solver.

3.2 The Planar-Pushing Markov Decision Process

The task is modelled as an episodic Markov Decision Process (MDP) with a macro-action, so that a single decision corresponds to one complete push rather than to a low-level joint command.

State. The state is a 28-dimensional vector composed of four groups. These are the end-effector pose (6 values), the object pose and velocity (14 values), the goal pose (6 values), and the two scalar position and orientation errors between object and goal. As explained in Section 2.1.2, the current observation does not report the underlying contact state. The problem is therefore formally partially observable, and every policy is recurrent.

Action. The policy does not command joint torques. It selects an *object-relative push primitive* described by four parameters: a radial approach offset r , an approach angle ϕ in the object frame, a push length ℓ , and a push direction θ in the world frame. Each parameter is discretised into 21 bins, so the action is a four-dimensional MultiCategorical variable with $21^4 = 194,481$ discrete values. Section 3.3 describes how the primitive is decoded into a robot motion and executed.

Transition. The object moves according to the quasi-static pushing mechanics of Section 2.8. The limit-surface normality condition of Equation 14 couples translation and rotation: almost every push changes both the position and the orientation of the block. This coupling is a property of the environment, not a modelling choice. It is the reason the task cannot be decomposed into independent position and orientation sub-problems.

Reward. The reward is the subject of Sections 3.4 to 3.6. It is the single component that differs across the models compared in this thesis, which is what makes the comparison a controlled one.

Success criterion. A push episode is counted as successful when the object is within 0.05 m of the goal position *and* within 0.2 rad of the goal orientation. This combined gate is applied identically in training and evaluation. The two conditions must hold at the same time, and the physics couples the two axes. The combined gate is therefore the central difficulty of the task, and the quantity every model is ultimately measured against.

3.3 The Object-Relative Push Primitive

Each decision commands one complete push rather than a joint torque, so the learning problem is defined over whole pushing motions. The four parameters are read in the object frame wherever possible, which is the property that makes the primitive robust to the object’s location. The radial approach offset r and the approach angle ϕ place the contact point on the object’s boundary. The angle ϕ selects a direction around the object, and r sets how far outside the boundary the arm first descends. The push length ℓ and the push direction θ then define the straight motion that follows contact. Because the contact point is expressed relative to the object, the arm meets the object wherever it lies on the table, and the contact rate stays near 95%.

This choice is deliberate, and it addresses two shortcomings of an absolute parameterisation in table coordinates. The first concerns exploration. When the push start was specified in world coordinates, a randomly sampled action contacted the object on only about 2% of pushes. The remaining pushes missed, returned no progress signal, and left the policy with almost nothing to learn from; in self-play the goal-proposing agent could not begin to improve. Placing the contact point on the object boundary instead makes contact the default, so nearly every push becomes an informative transition. The second concerns generalisation. With a fixed object position and absolute coordinates, the policy could memorise a single approach point and ignore the observation. Tying the action to the current object pose removes that shortcut, and requires the policy to read where the object is, so it acts on objects placed anywhere on the table. The push direction θ is kept in the world frame, aligned with the relative goal-direction feature of the observation, so the direction that reduces the goal distance is straightforward to learn. Each parameter is discretised into 21 bins, and the bins map linearly onto physical ranges bounded by the reachable workspace. Once decoded, the primitive runs as a five-phase Cartesian trajectory. The arm approaches a point above the contact location, descends to the table, pushes along θ for a distance ℓ , retracts, and returns to a neutral pose. The whole motion spans 72 physics substeps with the gripper held closed. Each Cartesian waypoint is mapped to arm joint targets by the cuRobo inverse-kinematics solver. A waypoint that falls outside the region where the solver returns a valid arm configuration is clamped to the boundary, so the credit assigned to an action matches the motion actually performed. Figure 2 shows the discrete choices the primitive exposes to the policy: the object-relative grid of contact points, and the world-frame push directions and lengths. Algorithm 1 summarises the decoding and execution; the workspace bounds and the five phases are detailed in Appendix E.




Figure 2: The object-relative push primitive, drawn to scale over a top-down T-block. Panel (a) shows the approach placement in the object frame: the contact point is set by a radial offset r , sampled from the band between 0.04 m and 0.12 m from the object centre (dashed circles), and an approach angle ϕ , one of 21 directions (grey spokes); the green points are the resulting grid of candidate contact points. The band is sized so that it encloses the object, whose outline reaches about 0.10 m from the tracked centre. Panel (b) shows the push itself in the world frame, from a representative contact point: the push direction θ is one of 21 headings (blue arrows), and the push length is one of 21 values from 0 to 0.20 m (marked along one heading). Each of the four parameters is discretised into 21 bins, so the action is a four-dimensional choice over a fixed grid rather than a continuous command.

Algorithm 1 Decode and execute one object-relative push primitive.

Require: action bins $(b_r, b_\phi, b_\ell, b_\theta) \in \{0, \dots, 20\}^4$; object pose (x_o, y_o, ψ_o)

- 1: map each bin to its physical value: offset r , angle ϕ , length ℓ , direction θ
 - 2: $p_{\text{contact}} \leftarrow$ boundary point of the object at angle ϕ , backed off by r
 - 3: $p_{\text{start}} \leftarrow$ world position of p_{contact}
 - 4: $p_{\text{end}} \leftarrow p_{\text{start}} + \ell (\cos \theta, \sin \theta)$
 - 5: clamp p_{start} and p_{end} to the reachable workspace
 - 6: **for** phase in {approach, descend, push, retract, return} **do**
 - 7: compute the Cartesian waypoint for the phase
 - 8: solve arm joint targets for the waypoint with cuRobo IK
 - 9: step the simulator with the gripper held closed
 - 10: **end for**
 - 11: **return** the object pose after the push
-

3.4 Why the Improvement Reward Underperforms

An initial baseline used a hand-tuned reward based on fractional improvement. It rewards each push in proportion to the fraction of the remaining error it removes. Let d_{prev} and d_{now} be the position error before and after a push, and y_{prev} , y_{now} the orientation error. The reward is then

$$R = \alpha \frac{d_{\text{prev}} - d_{\text{now}}}{d_{\text{prev}}} - \beta_{\text{pos}} d_{\text{now}} - \beta_{\text{rot}} y_{\text{now}}, \quad \alpha = 3.0, \beta_{\text{pos}} = 0.5, \beta_{\text{rot}} = 0.25. \quad (16)$$

This reward performs poorly for two related reasons. Both concern the gradient it supplies to PPO, not the optimal policy it defines. Consider a representative push that moves the object 1 cm closer to a goal 30 cm away. The improvement term contributes $3.0 \times 0.01/0.30 \approx 0.10$. The two penalty terms contribute $-0.5 \times 0.29 \approx -0.145$ and $-0.25 \times 0.5 = -0.125$, for a net reward of about -0.17 . A push that makes genuine progress therefore receives negative reward, because the standing penalty on the remaining distance dominates the small improvement. This is the first problem: the reward is negative for most useful pushes, so the policy gradient points away from goal-directed behaviour.

The second problem is near-zero variance. Across a rollout, roughly 99% of pushes produce a reward in the narrow interval $[-0.5, +0.1]$, and the sparse completion bonuses fire on about 1% of pushes. At 528 environments and 15 pushes per rollout, only about 79 bonus events occur per iteration, too few to overcome the noise in the dense term. The Generalized Advantage Estimation (GAE) advantage is then dominated by noise, and the gradient is uninformative. The failure is therefore one of gradient quality and reward variance, not of a mis-specified optimum. The reward defines a reasonable objective, yet it supplies a poor learning signal.

3.5 Potential-Based Reward Shaping for Pushing

The remedy is to replace the improvement reward with a potential-based reward. The per-push signal is then dense, correctly signed, and bounded, while the optimal policy is provably unchanged (Section 2.4). The design uses two bounded exponential potentials, one for position and one for orientation:

$$\Phi_{\text{pos}}(s) = w_{\text{pos}} \exp(-k_p d^2), \quad (17)$$

$$\Phi_{\text{rot}}(s) = w_{\text{rot}} \exp(-k_r c(s)), \quad c(s) = \frac{1 - \cos(\theta^* - \theta)}{2}, \quad (18)$$

where d is the planar position error and $c(s)$ is a cosine orientation distance. The cosine form is smooth and free of the discontinuity a modular angle difference would introduce at $\pm\pi$; θ^* and θ are the goal and object yaw. The parameters are $k_p = 30$, $k_r = 5$, and $w_{\text{pos}} = w_{\text{rot}} = 10$. The dense reward on a transition is the difference of the summed potential,

$$F(s, s') = [\Phi_{\text{pos}}(s') - \Phi_{\text{pos}}(s)] + [\Phi_{\text{rot}}(s') - \Phi_{\text{rot}}(s)], \quad (19)$$

which is Equation 10 with $\gamma_{\text{shaping}} = 1$. The undiscounted setting is admissible because the task is episodic and every episode terminates, as established for episodic MDPs by [19]. Two sparse bonuses are layered on the dense term. A completion bonus of +5 is given when the position gate is satisfied, and a further +2 when the combined gate is satisfied. A penalty of -5 is applied if the block tips over. These bonuses are not themselves potential-based, but they fire only at terminal or near-terminal events, where they do not create the reward-collecting cycles that Section 2.4 identifies as the danger of non-potential shaping.

On the same representative push, the potential-based reward inverts the sign of the baseline. For a 1 cm improvement from 30 cm, the position potential difference is $10 \times [\exp(-30 \times 0.29^2) - \exp(-30 \times 0.30^2)] \approx +0.13$. Every push toward the goal now receives positive reward, and every push away from it negative reward. Because the potential is bounded in $[0, w]$, the reward magnitude cannot be inflated by a physics contact spike. Figure 3 contrasts the per-push reward of the two schemes over a representative episode, and shows the potential landscape and the gradient comparison.




Figure 3: Left: per-push reward over a representative five-push episode for the improvement reward and the potential-based reward. The improvement reward is negative or near-zero for most pushes; the potential-based reward is correctly signed on every push. Right: the gradient supplied by the two schemes as a function of distance to the goal.

The reward is implemented as a potential-based reward function. Four properties make it suitable for a contact-rich task, and all follow from its form. It is Markovian, since it depends only on the current and successor state. It is bounded. It produces zero net reward over any cycle that returns to the same state, which removes the incentive for reward-collecting loops. And it produces zero reward while the object is stationary, so it imposes no penalty for deliberation.

3.6 The SE(2) Unified Distance

The two-potential reward treats position and orientation as separate terms whose gradients can point in different directions. For the self-play models (Section 3.7) the reward instead uses a single SE(2) distance, so one potential governs the whole pose. Following the Lie-group formulation of Section 2.8, the pose error is summarised by the scalar

$$d_{\text{pose}} = \sqrt{\Delta x^2 + \Delta y^2 + L^2 \Delta \theta^2}, \quad (20)$$

where $\Delta x, \Delta y$ are the position error components and $\Delta \theta$ the orientation error. Here L is a characteristic length that converts an angular error into an equivalent linear displacement. For the T-block $L = 0.07$ m, the block’s radius of gyration; it is a physical property of the object, not a tuning parameter. The corresponding single potential and its dense reward are

$$\Phi_{\text{pose}}(s) = w \exp(-k d_{\text{pose}}^2), \quad F(s, s') = \Phi_{\text{pose}}(s') - \Phi_{\text{pose}}(s), \quad (21)$$

with $k = 30$ and $w = 10$. The success threshold for this metric is $d_{\text{pose}} < 0.055$ m for the T-block. The single potential removes the competing-gradient problem of the two-potential form.

A push that reduces d_{pose} is rewarded whether the reduction comes from position or orientation, so the reward is aligned with the coupled physics of the task. Both the two-potential and the unified formulations share the same reward implementation.

3.7 The Models

All nine models share the same backbone: a single PPO agent with an LSTM of hidden size 256 followed by a multilayer perceptron, and a MultiCategorical head over the four-dimensional, 21-bin push primitive. They differ only in the reward and the training dynamics, which is what makes the comparison controlled. The models fall into four groups by the role each serves. The single-agent anchors establish that the goal-reaching task is solvable and measure the reward-efficiency result. The self-play subjects are the primary objects of study: the two variants that underperform the anchors on the coupled goal. The coupling-isolation models repeat self-play on a rotationally symmetric disc, whose position-only goal removes the coupled success criterion. The reward-structure control and the cross-environment counterparts test two further explanations for the self-play result. Table 1 lists all nine, the number of random seeds trained per model, and what each is meant to show.

Model	Group	Reward and training dynamics	Envs	Seeds	What it is meant to show
PPO-Baseline	Single-agent anchor	Improvement reward (Equation 16); one agent	256–2,048	5	The goal is reachable without shaping; sets the reference the potential-based reward is measured against.
PPO-PBRS	Single-agent anchor	Two-potential PBRS over position and orientation; one agent	256–2,048	5	The solvable anchor; reaches success with far fewer environments than the improvement reward.
PPO-Curriculum	Single-agent anchor	PBRS with forced position-then-orientation staging	528	5	Whether staging the two objectives helps once the reward is well shaped.
ASP-dPose	Self-play subject	SE(2) PBRS for Bob; outcome-based Alice reward; imitation on	256, 528, 2,048	5	Primary subject: outcome-reward self-play on the coupled T-block goal; trained up to 2,048 environments to test compute scale.
TASP-dPose	Self-play subject	SE(2) PBRS for Bob; time-based Alice reward; imitation on	256, 528, 2,048	5	Primary subject; isolates Alice’s reward against ASP-dPose, and is likewise scaled to 2,048 environments.
ASP-disc	Coupling isolation	Outcome-based self-play on a disc; position-only goal	528	5	Whether self-play reaches the anchor level once orientation is unconstrained.
TASP-disc	Coupling isolation	Time-based self-play on a disc; position-only goal	528	5	Time-reward counterpart of the disc isolation.
Bob-penalty	Reward-structure control	Symmetric self-play reward; Bob penalised for its own solve time	528	5	Whether a symmetric reward structure collapses the self-play loop.
gym A/B/C	Cross-environment	Single agent, curriculum, and self-play in the two-dimensional simulator	64, 256	1	Whether the ordering of the methods holds in a second simulator.

Table 1: The nine models. Every model uses the same PPO and LSTM backbone and the same object-relative push primitive; only the reward and the training dynamics change. The seed column gives the number of random seeds behind each model’s headline confidence interval, taken at 528 environments. The single-agent Baseline and PBRS are trained across 256, 512, 1,024, 2,048, and 528 environments, with three seeds per swept count and five at the 528 anchor. The self-play subjects ASP-dPose and TASP-dPose are additionally trained at 256 and 2,048 environments, with three seeds, so that the single-GPU compute budget can be tested, rather

PPO-Baseline. A single PPO agent trained with the improvement reward of Equation 16. It serves as the reference point for the reward comparison of Section 3.4 and is trained at 2,048 parallel environments, a larger batch than the other models use.

PPO-PBRS. A single PPO agent trained with the two-potential reward of Equations 17 to 19. This is the simplest model that uses the potential-based reward, with no curriculum and no second agent, and it is trained at 528 environments.

Algorithm 2 states the single-agent training loop shared by PPO-Baseline and PPO-PBRS; the two differ only in the reward R .

Algorithm 2 Single-agent training (PPO-Baseline and PPO-PBRS).

Require: reward R ; environments N ; iterations T ; pushes per rollout H ; epochs K

```

1: initialise recurrent policy  $\pi_\theta$  and value function  $V_\varphi$ 
2: for iteration = 1 to  $T$  do
3:   for each of the  $N$  environments in parallel do
4:     sample an object start pose and a goal pose
5:     for  $t = 1$  to  $H$  do
6:       observe state  $s_t$ ; sample push  $a_t \sim \pi_\theta(\cdot \mid s_t)$ 
7:       execute  $a_t$  (Algorithm 1); observe  $s_{t+1}$ 
8:        $r_t \leftarrow R(s_t, s_{t+1})$ 
9:     end for
10:  end for
11:  estimate advantages with Generalized Advantage Estimation
12:  for epoch = 1 to  $K$  do
13:    update  $\pi_\theta$  and  $V_\varphi$  with the clipped PPO objective
14:  end for
15: end for
16: return  $\pi_\theta$ 

```

PPO-Curriculum. PPO-PBRS with a forced three-phase curriculum that stages the two objectives. Phase 1 trains position only, with the orientation weight set to zero. Phase 2 begins when the episodic position-only success rate reaches 0.5 over a 20-iteration window; the orientation weight is then increased from 0 to 10 over 200 iterations. Phase 3 is the full two-objective reward. The trigger gates on the episodic position-only success rate, rather than on a per-push error mean. This is the corrected version of a trigger that in an earlier implementation never activated; the earlier failure and its correction are documented in Appendix H. Algorithm 3 shows the staging around the single-agent loop.

Algorithm 3 Forced curriculum staging (PPO-Curriculum).

Require: potential-based reward with orientation weight w_{rot} ; iterations T

```

1:  $w_{\text{rot}} \leftarrow 0$ ; phase  $\leftarrow 1$  ▷ Phase 1: position only
2: for iteration = 1 to  $T$  do
3:   run one iteration of Algorithm 2 with the current  $w_{\text{rot}}$ 
4:   if phase = 1 and the position-only success rate  $\geq 0.5$  over the last 20 iterations then
5:     phase  $\leftarrow 2$  ▷ Phase 2: raise the orientation weight
6:   end if
7:   if phase = 2 then
8:      $w_{\text{rot}} \leftarrow \min(10, w_{\text{rot}} + 10/200)$ 
9:     if  $w_{\text{rot}} = 10$  then
10:      phase  $\leftarrow 3$  ▷ Phase 3: full two-objective reward
11:    end if
12:  end if
13: end for
14: return  $\pi_\theta$ 
```

ASP-dPose. An asymmetric self-play model that uses the SE(2) reward of Equation 21 for the goal-solving agent, Bob. The loop is shown in Figure 4. The goal-setting agent, Alice, performs five pushes to propose a goal. The goal is accepted only if it passes a validity check: the block has moved sufficiently, remains on the table, and lies inside the workspace. Bob then attempts the goal from a clean reset over up to ten pushes. Alice’s reward is *outcome-based*, following [9]: she receives a reward when Bob fails, a shaping bonus for a valid goal, and a penalty for an out-of-bounds goal. Bob additionally uses a GoalEncoder, a small multilayer perceptron that maps the six-dimensional goal pose into an eight-dimensional latent, in the difference variant of [35]. Bob is trained with Alice Behavioral Cloning at $\beta = 0.5$. This imitation term is enabled in the canonical model; a variant with it disabled is retained only as an ablation.

TASP-dPose. Identical to ASP-dPose except that Alice’s reward is *time-based*, following [8]: $R_A = \gamma_{sp} \max(0, t_B - t_A)$. This rewards Alice for proposing goals that take Bob more pushes to solve than Alice took to set, subject to a penalty on Alice’s own effort. It is the self-regulating reward of Equation 12. The two self-play models therefore differ only in Alice’s reward, which isolates the effect of the goal proposal mechanism.




Figure 4: The asymmetric self-play loop. Alice proposes a goal by acting for a fixed number of pushes; the goal is validated; Bob then attempts it from a clean reset. Alice’s reward differs between ASP-dPose (outcome-based) and TASP-dPose (time-based); Bob’s SE(2) potential-based reward is the same in both.

Algorithm 4 states the self-play training loop shared by both variants; the disc and Bob-penalty models are the marked variations of it.

Algorithm 4 Asymmetric self-play training (ASP-dPose and TASP-dPose).

Require: Bob reward R_B (SE(2) potential-based); Alice reward mode; imitation weight β ;
 Alice pushes H_A ; Bob pushes H_B

- 1: initialise goal-setting policy π_A , goal-solving policy π_B with its goal encoder, and the value functions
- 2: **for** iteration = 1 to T **do**
- 3: **for** each environment in parallel **do**
- 4: reset the scene; Alice acts for H_A pushes and records trajectory τ_A
- 5: $g \leftarrow$ object pose after Alice’s pushes $\triangleright$ proposed goal
- 6: **if** g is invalid (off the table or outside the workspace) **then**
- 7: penalise Alice; skip Bob
- 8: **else**
- 9: reset the scene to the same start; Bob attempts g for up to H_B pushes, sampling $a \sim \pi_B(\cdot \mid s, g)$
- 10: $t_B \leftarrow$ Bob’s push count; record whether the combined gate is reached
- 11: **if** Alice reward is outcome-based **then**
- 12: $R_A \leftarrow$ reward when Bob fails, plus a valid-goal shaping bonus
- 13: **else**
- 14: $R_A \leftarrow \gamma_{sp} \max(0, t_B - t_A)$ $\triangleright$ time-based
- 15: **end if**
- 16: $\triangleright$ Bob-penalty variant additionally sets $R_B \leftarrow R_B - \gamma_{sp} t_B$
- 17: **end if**
- 18: **end for**
- 19: update π_A with PPO on R_A
- 20: update π_B with $\mathcal{L} = \mathcal{L}_{RL} + \beta \mathcal{L}_{abc}$, where $\mathcal{L}_{abc}$ clones Alice’s actions on failed goals
- 21: **end for**
- 22: **return** π_B

ASP-disc and TASP-disc. Two self-play models identical to ASP-dPose and TASP-dPose, except that the manipulated object is a rotationally symmetric disc rather than the T-block. A disc has no distinguished orientation, so its goal constrains position alone, and the combined success gate reduces to a position-only criterion. The mechanical reason a disc removes the orientation coupling that the T-block enforces is set out in Section 2.8 and Appendix A. These two models isolate the coupled multi-objective gate as a single variable. If self-play reaches the single-agent level on the disc yet underperforms it on the T-block, the coupled gate, rather than the reward or the goal encoder, is the factor that separates the two.

Bob-penalty. A symmetric variant of time-based self-play in which the goal-solving agent is penalised for its own solve time, in addition to the goal-proposing agent being rewarded for it. It tests whether the asymmetry of the reward structure is necessary in this domain. The behaviour of this model, a fail-fast equilibrium in which immediate failure produces a higher return than a longer attempt, is analysed in Appendix I.

Cross-environment counterparts. The single-agent, curriculum, and self-play models are retrained in a second, two-dimensional pushing simulator that uses neither inverse kinematics nor three-dimensional contact physics. These counterparts run on the central processing unit

at environment counts of 64 and 256. Their role is to test whether the ordering of the methods holds in an environment of different fidelity, so that a conclusion does not rest on one simulator. The protocol is described in Section 4.4.

3.8 Why PBRs Assists Self-Play but Does Not Solve It

Potential-based shaping improves self-play but does not make it competitive with single-agent training, and the reason is visible in the structure of the GAE advantage. The advantage estimate is a discounted sum of one-step temporal differences, each of which decomposes into a reward term and a value-bootstrap term:

$$\hat{A}_t = \sum_{l \geq 0} (\gamma \lambda)^l \underbrace{[\Phi(s_{t+l+1}) - \Phi(s_{t+l})]}_{\text{potential term, correctly signed}} + \sum_{l \geq 0} (\gamma \lambda)^l \underbrace{[\gamma V(s_{t+l+1}) - V(s_{t+l})]}_{\text{value term, biased when } V \text{ is stale}}. \quad (22)$$

The potential term is correctly signed on every push, regardless of the accuracy of the value function, because the potential increases whenever the object moves toward the goal. The value term, by contrast, depends on the accuracy of V . In single-agent training the goal distribution is stationary, so V is accurate and both terms agree. Under self-play, Alice changes the goal distribution at every iteration, so V is perpetually estimated on a distribution that no longer holds, and the value term is biased. The potential term supplies a floor of correct learning signal, which lifts self-play above the level it reaches under the sparse or improvement reward. It cannot, however, remove the bias contributed by the stale value term. The consequence, quantified in Chapter 5, is that **potential-based shaping raises self-play from near-zero to a modest success rate without approaching the single-agent result.**

This account is a hypothesis about the learning dynamics, not a result. Both of its claims, the elevated value error under a shifting goal distribution and the suppression of the imitation signal, are measured directly through the training instrumentation described in Section 3.11, and the measurements are reported in Chapter 5 rather than asserted here.

3.9 Experiment-Set Overview

The nine models of Section 3.7 are the objects of study. A single training run of each, at one batch size and one seed, cannot support the statistical and causal claims this thesis makes. The evaluation is therefore organised as an experiment set of budget-matched training runs, summarised in Table 2. Its per-phase configurations and job counts are given in Chapter 4. This section states the two principles that make the experiment set a controlled study, together with the validation protocol and training instrumentation that every run shares. The numerical outcomes appear in Chapter 5.

The first principle is **multiple seeds with reported confidence intervals**. Single-seed reinforcement-learning results are known to be unreliable, because the variance between random seeds can exceed the difference between methods [53, 54]. Every headline comparison in this thesis is therefore repeated across several seeds. Each is reported as a mean with a 95% confidence interval, so a claimed difference between two models is distinguished from seed noise.

The second principle is **budget matching**. The aim is to compare batch sizes without confounding them with the amount of experience. The total number of pushes seen over training is therefore held constant while the number of parallel environments is varied. The iteration

count is set inversely proportional to the number of environments, so each configuration consumes the same total experience of about 23.8 million pushes. The temporal window of the recurrent policy is fixed at 15 pushes per rollout; only the number of parallel environments, and therefore the per-update batch size, changes. This isolates the per-update batch size, the quantity that drives the value-function bias of Section 3.8, from the total experience. Budget matching applies to the Isaac Lab runs; the second-simulator runs instead use a fixed iteration count, as described in Chapter 4. To examine whether the single-GPU budget is itself the cause of the self-play result, the two self-play subjects are additionally trained at 256 and 2,048 environments, so that compute scale is tested rather than assumed.

Two conventions hold across every self-play run in the experiment set. The imitation term is enabled in the canonical self-play models, and disabled only for the ablation arm. Both the single-agent and the self-play validators sample actions from the policy, rather than taking the most probable action, so the two families are compared under one protocol.

Phase	Question	Design
Anchor and efficiency	Is the goal-reaching task solvable, and how few environments does the potential-based reward need?	PPO-Baseline and PPO-PBRS across four environment counts, three seeds each
Seed intervals	Is each model difference larger than seed noise?	the baseline, the potential-based agent, the curriculum, and both self-play subjects at the anchor budget, five seeds
Coupling isolation	Does self-play reach the anchor level when the goal is position-only?	the two disc self-play models against the two T-block subjects, five seeds
Self-play scale	Does more compute rescue self-play, or is scale not the cause?	both self-play subjects at 256 and 2,048 environments, three seeds
Component ablations	Which component of the self-play package matters?	the self-play subjects with the imitation term, then the goal encoder, removed one at a time, three seeds
Imitation positive control	Does imitation help when the task is easy?	disc self-play with the imitation term enabled and disabled, three seeds
Reward structure	Does the symmetric reward collapse the loop?	the Bob-penalty model at the anchor budget, five seeds
Cross-environment	Does the ordering hold in a second simulator?	single agent, curriculum, and self-play in the two-dimensional simulator across environment counts 64 and 256 and push lengths 15 and 45, single seed

Table 2: The experiment set. Each phase answers one question; the per-phase training configurations, environment counts, seeds, and job counts are tabulated in Chapter 4.

3.10 Validation Protocol

Every trained policy is evaluated on the same held-out suite of 30 T-block scenes, illustrated in Figure 5. The suite is divided into three groups of ten. The first ten are rotation-dominant scenes with large orientation changes. The next ten are position-only scenes whose goal orientation equals the start orientation. The last ten are combined scenes that require both a position and an orientation change. Within each group the difficulty ranges from easy to hard, graded by the start-to-goal distance and the required rotation. The held-out scenes are fixed in advance and never seen during training, so performance on them measures generalisation rather than memorisation.




Figure 5: The 30-scene held-out validation suite. Each panel shows the block start pose and the goal pose for one scene. The suite is identical for every model and every seed.

Each scene is attempted 20 times, and within each attempt the policy is allowed up to 30 pushes. Two success measures are reported and kept distinct throughout the thesis. The **scene success rate** counts a scene as solved if at least one of its 20 attempts satisfies the combined gate of Section 3.2. This is the headline measure, and follows the convention common in manipulation evaluation, where a task counts as achievable if the policy can achieve it at all. The **trial success rate** is the fraction of individual attempts that satisfy the gate, and is always the lower of the two. Reporting both prevents the best-of-20 protocol from overstating performance, and the gap between them is stated wherever the headline number appears.

Three protocol decisions make the numbers comparable across models. Each corrects a subtle way the comparison could otherwise be invalidated. First, the checkpoint selected for validation is the one that maximised the *combined* position-and-orientation success rate during training, not the looser single-distance threshold used internally by the self-play reward. Selecting on a different criterion than the one reported would flatter the self-play models. Second, the self-play policies trained on the rotationally symmetric disc are validated on disc scenes, whose goal orientation is unconstrained, not on the T-block scenes they never trained on. Validating an agent on a task it never trained for would void its result. Third, checkpoints that lack a combined-gate best model, and are validated only from their latest checkpoint, are recorded separately and excluded from the confidence-interval aggregation. A partially-trained run therefore cannot contaminate a headline mean. The per-seed results are aggregated into a mean with a 95% confidence interval per model, batch size, and iteration count.

3.11 Training Instrumentation for the Self-Play Mechanism

The account in Section 3.8 of why potential-based shaping only partially assists self-play rests on two claims. The first is that the value-function term of the advantage estimate is biased under the shifting goal distribution. The second is that the behavioural-cloning signal is suppressed when the goal-proposing agent’s demonstrated actions lie far from the goal-solving agent’s policy. Rather than leave these as arguments, the training runs are instrumented so that both are measured directly, and the measurements are reported in Chapter 5.

For the value-function claim, the mean absolute difference between the value estimates and the empirical returns is logged at every policy update. This is done for the single-agent, self-play, and time-based self-play models alike. Under the budget-matched scale sweep, this quantity can be plotted against the number of parallel environments. The mechanism of Section 3.8 makes two predictions: the value error is larger under self-play than under single-agent training at the same batch size, and it falls as the batch grows. A value error that declines toward the single-agent level with scale would support a budget-limitation account. One that remains elevated would indicate a structural deficit independent of budget.

For the behavioural-cloning claim, the imitation term is instrumented in three ways: the share of demonstrated actions whose probability ratio falls outside the clipping interval, the mean ratio, and the mean probability the goal-solving agent assigns to the goal-proposing agent’s demonstrated actions. The clipping mechanism inherited from the imitation loss of [9] zeroes the gradient for any demonstrated action whose ratio leaves the trust interval. A large clipped fraction is therefore the direct signature of a suppressed imitation signal. Measuring this fraction converts the clip-saturation account into a reported quantity. A positive control on the rotationally symmetric disc, where behavioural cloning is enabled and the task is easy, tests whether the mechanism can help at all in this setting. If imitation assists anywhere it should assist there, so its failure on the coupled T-block task cannot then be dismissed as a broken implementation.

4 Experimental Setup

This chapter describes how the nine models of Chapter 3 are trained and evaluated. It specifies the tools and hardware, the training configuration and the experiment set, and the held-out validation protocol. It then describes the cross-environment comparison against a second simulator, an off-policy baseline based on Soft Actor-Critic, and the metrics reported in the results.

4.1 Tools and Infrastructure

Training and Isaac-Lab evaluation run on a single RTX Pro 6000 GPU with 96 GB of memory. The software stack is Isaac Sim 5.1.0, Isaac Lab 2.3.0, cuRobo 0.7.5 for batched Inverse Kinematics (IK), and PyTorch 2.7.0. A second environment, the two-dimensional `gym-pusht` simulator, is used for the cross-environment comparison of Section 4.4. It runs on 32 CPU cores and uses neither IK nor three-dimensional contact physics. The off-policy baseline of Section 4.5 uses Stable-Baselines3 2.7.0. Running the same models in two simulators of different fidelity separates conclusions that hold across environments from artefacts of a single implementation.

4.2 Training Configuration and the Experiment Set

At the 528-environment anchor, every model is trained for approximately 3,000 iterations. This corresponds to about 1.6 days of wall-clock time per model on the RTX Pro 6000. At 15 pushes per rollout it yields a batch of roughly 7,920 push transitions per PPO update. For the reward-efficiency comparison, the improvement-reward baseline and the potential-based agent are trained across the same environment counts, from 256 to 2,048, so the two rewards are measured at matched parallelism rather than at a single batch. The two self-play subjects are additionally trained at 256 and 2,048 environments, which tests whether the single-GPU compute budget is the cause of their result. The PPO hyperparameters are a learning rate of 3×10^{-4} , a clipping range of 0.2, a discount factor of 0.998, a GAE parameter of 0.95, and an entropy coefficient of 0.005. The self-play models additionally use Alice Behavioral Cloning at $\beta = 0.5$.

Two points about the training history are stated for reproducibility. First, all reported checkpoints post-date a sequence of corrections to the reward and the environment. The reported numbers are therefore produced on the corrected code, and earlier runs are excluded. Second, the improvement-reward baseline and the potential-based agent are compared across a shared range of environment counts, so that any difference between the two rewards cannot be attributed to unequal parallelism. The budgets are tabulated in Appendix K.

Every Isaac Lab run in the experiment set is budget-matched. The total experience is held at about 23.8 million pushes, and the iteration count is set inversely proportional to the number of parallel environments, so a change in the number of environments changes the per-update batch size alone. This design has one confound that is stated wherever a scale result appears: a larger batch also means fewer gradient updates over training, so batch size and update count move together and cannot be fully separated. Table 3 lists the phases, the models each one trains, the environment counts, the matched iteration counts, and the seeds.

Phase	Models	Envs	Iterations	Seeds	Runs
Anchor and efficiency	PPO-Baseline, PPO-PBRs	256–2,048	773–6,188	3	24
Seed intervals	Baseline, PBRs, Curriculum, ASP-dPose, TASP-dPose	528	3,000	5	25
Coupling isolation	ASP-disc, TASP-disc	528	3,000	5	10
Self-play scale	ASP-dPose, TASP-dPose	256, 2,048	6,188; 773	3	12
Component ablations	ASP-dPose, TASP-dPose	528	3,000	3	12
Imitation positive control	ASP-disc	528	3,000	3	6
Reward structure	Bob-penalty	528	3,000	5	5
Cross-environment	single agent, curriculum, self-play	64, 256	3,000	1	10

Table 3: The phases of the experiment set. Environments and iterations trade off to hold the total experience at about 23.8 million pushes for the Isaac Lab runs; the cross-environment runs instead use a fixed 3,000 iterations and a single seed, and vary the per-update batch through the environment count and the push length. The experiment set comprises 94 training runs in Isaac Lab and 10 in the second simulator.

4.3 Validation Protocol

Every model is evaluated on the same held-out suite of 30 T-block scenes, illustrated in Figure 6. The suite is divided into three groups of ten. The first ten are rotation-dominant scenes with large orientation changes. The next ten are position-only scenes whose goal orientation equals the start orientation. The last ten are combined scenes that require both a position and an orientation change. Within each group the difficulty ranges from easy to hard, graded by the start-to-goal distance and the magnitude of the required rotation.




Figure 6: The 30-scene held-out validation suite. Each panel shows the block start pose and the goal pose for one scene. The suite is identical for every model.

Each scene is attempted 20 times, and within each attempt the policy is allowed up to 30 pushes. Both the single-agent and the self-play validators sample actions from the policy, rather than taking the most probable action, so every model is scored under one protocol. Two success measures are reported and kept distinct throughout. The *scene success rate* counts a scene as solved if at least one of its 20 attempts satisfies the combined gate of Section 3.2. This is the

headline number, and follows the convention common in manipulation evaluation. The *trial success rate* is the fraction of individual attempts that satisfy the gate, and is always lower than the scene success rate. Reporting both prevents the best-of-20 protocol from overstating performance, and the gap between them is stated wherever the headline number appears; the final values are **TBD** until the seed runs are collated. Each model is evaluated across its trained seeds, and the results are aggregated into a mean with a 95 % confidence interval per model, batch size, and iteration count.

4.4 Cross-Environment Protocol

To test whether the ordering of the models is a property of the methods rather than of Isaac Lab, the single-agent, curriculum, and self-play models are retrained and evaluated in `gym-pusht` under the same 30-scene gate. The two environments differ mainly in the size of the PPO batch. Isaac Lab produces about 7,920 transitions per update at the 528-environment anchor, while `gym-pusht` produces between about 960 and 11,520, depending on the environment count and the push length. The gym runs use a fixed 3,000 iterations and sweep the per-update batch through these two settings, so the comparison tests whether the ordering of the methods survives a second, lower-fidelity simulator, rather than matching the total experience between the two environments. Figure 7 summarises the throughput and batch size of the two environments.




Figure 7: Throughput and PPO batch size for Isaac Lab and `gym-pusht`. Raw throughput is comparable; the per-update batch differs by roughly a factor of eight, which is the variable the cross-environment comparison isolates.

4.5 Off-Policy Baseline: SAC with Hindsight Experience Replay

The principal models are all on-policy. To provide an external point of calibration from a different algorithm family, an off-policy baseline is built from Soft Actor-Critic (SAC) [16] combined with Hindsight Experience Replay (HER) [18]. SAC is a natural counterpart to PPO for this problem. It is sample-efficient through experience replay, and its maximum-entropy objective (Equation 9) supports exploration under sparse reward. HER addresses the sparse goal-conditioned setting directly: it relabels the object pose actually reached in a failed attempt as if it had been the goal, so every attempt yields usable transitions.

The baseline runs on a direct control environment through Stable-Baselines3, and reuses the shared push primitive and the cuRobo IK pipeline of Section 3.2. Two adaptations are required. First, SAC is a continuous-action algorithm, so the four push parameters become a continuous action in $[-1, 1]^4$ and are decoded to the same object-relative primitive, rather than

the MultiCategorical variable used by the PPO models. Second, the goal-conditioned observation is expressed in the dictionary form the replay buffer expects, with separate observation, achieved-goal, and desired-goal entries. The baseline is trained and evaluated under the same 30-scene protocol as the principal models, so its success rate is directly comparable. Its role is calibration rather than competition. It indicates whether an off-policy, replay-based method reaches performance similar to the on-policy potential-based model, and its status against the compute budget is reported with the results.

4.6 Metrics

The results report the following quantities, each with its sample size and protocol stated at the point of use. The primary metric is the scene success rate over the 30-scene suite, with the trial success rate reported alongside it. The breakdown by scene type, position-only against combined, is reported because it localises where each model fails. Mean position error and mean orientation error are reported as continuous measures of near-misses. For the self-play models the training logs also provide the per-axis success rates: the fraction of pushes that satisfy the position gate alone and the orientation gate alone. These are compared against the combined-gate rate to diagnose whether the two skills are learned but not combined. Average push count per attempt is reported as a measure of efficiency. Three further quantities are logged during self-play training and reported in the results. The first is the mean absolute value error, which measures the bias of the value function under the shifting goal distribution. The second is the share of demonstrated actions the imitation term clips away, a direct measure of how strongly the imitation signal is suppressed. The third is the Alice-replay rate: the fraction of proposed goals that the goal-proposing agent’s own action sequence reaches when replayed from the goal-solving agent’s start state. All reported numbers trace to the validation records summarised in the results.

5 Results

OLD RESULTS — pending experiment set: the entire chapter still reports the old single-seed numbers and the old research-question framing; every table, figure, and quoted rate below is provisional and will be regenerated from the experiment set (seed means with 95% confidence intervals), then reorganised around the new RQ1–RQ3.

This chapter reports the outcome of the five models on the 30-scene validation suite, organised by research question. Section 5.1 addresses the environment-efficiency of potential-based shaping (RQ1), Section 5.2 compares curriculum and self-play against single-agent training (RQ2), and Section 5.3 examines whether self-play learns the two constituent skills (RQ3). Section 5.4 confirms the ordering in a second simulator, and Section 5.5 reports qualitative behaviour and the off-policy baseline. Unless stated otherwise, every number is for the 30-scene suite, best of 20 attempts, up to 30 pushes, under the combined gate of Section 3.2. Table 4 collects the principal numbers.

Model	Scene SR	Trial SR	Pos-only	Pos+rot	Envs
PPO-Baseline (improvement reward)	80.0 %	70.5 %	100 %	35 %	2,048
PPO-PBRS	80.0 %	66.0 %	100 %	70 %	528
PPO-Curriculum	76.7 %	68.3 %	100 %	65 %	528
TASP-dPose (time-based self-play)	16.7 %	1.2 %	30 %	10 %	528
ASP-dPose (outcome-based self-play)	6.7 %	0.3 %	20 %	0 %	528

Table 4: Principal validation results on the 30-scene held-out suite (best of 20 attempts, up to 30 pushes, combined gate). Scene SR counts a scene solved if any of its 20 attempts passes the gate; trial SR is the per-attempt rate. The PPO-Baseline row uses the improvement reward and the original, non-simplified potential formulation reported at the position-only gate; its combined-gate performance is discussed in the text.

5.1 RQ1: Environment-Efficiency of Reward Shaping

The potential-based reward reaches the same scene success rate as the improvement reward while using a quarter of the parallel environments. PPO-PBRS attains a scene success rate of 80.0 % at 528 environments, matching the 80.0 % reached by the improvement-reward baseline at 2,048 environments. The reduction from 2,048 to 528 is a factor of roughly four in the parallelism required for the same performance. This answers RQ1 in the affirmative: for this task, potential-based shaping reduces the number of parallel environments needed for on-policy learning.

The benefit of shaping is larger on the combined objective than the headline number alone conveys. On the ten position-only scenes both rewards reach 100 %, so the parallelism reduction is the whole of the effect there. The ten combined scenes tell a different story. There, the potential-based reward reaches 70 %, while the earlier non-simplified formulation evaluated at the same gate reaches only 35 %, an improvement of 35 percentage points. The per-scene error plot for PPO-PBRS in Figure 8 shows why the six failing scenes fail. All six are rotation-

dominant: position is solved on every scene, and the residual failures are limited by orientation error rather than position error. This localises the remaining difficulty of the single-agent task to the orientation component of the combined gate.




Figure 8: Per-scene position and orientation error for PPO-PBRS on the 30-scene suite. Position is within tolerance on every scene; the failing scenes are limited by orientation error on rotation-dominant scenes.

5.2 RQ2: Curriculum and Self-Play Against Single-Agent

Neither the forced curriculum nor either self-play variant improves on the plain single-agent model. PPO-Curriculum reaches 76.7%, which is 3.3 percentage points below PPO-PBRS at 80.0%. Across the 30 independent scenes this difference is not statistically significant: the standard error of the difference is approximately ± 10.6 percentage points, so the two models are equivalent on this evidence. The staged curriculum does change the error profile. It attains a lower mean position error (0.023 m against 0.032 m), consistent with its position-only first phase, but a slightly higher mean orientation error (0.663 rad against 0.568 rad). The staging therefore reallocates precision between the two axes without improving the combined result.

The self-play models perform far worse than single-agent training. Figure 9 shows the full comparison. TASP-dPose, the stronger variant, reaches 16.7%, and ASP-dPose reaches 6.7%. These are factors of roughly 4.8 and 11.9 below PPO-PBRS. The time-based Alice reward of TASP-dPose is markedly better than the outcome-based reward of ASP-dPose, a difference of 10 percentage points. This indicates that the self-regulating time-based formulation of Section 2.6 produces a more usable goal distribution than the outcome-based one, yet neither approaches the single-agent result. A controlled measurement puts the per-iteration cost of self-play at roughly equal to that of single-agent training. The gap therefore cannot be attributed to self-play receiving less compute: it is a structural difference between the two training regimes, not a budget artefact.




Figure 9: Scene success rate (left) and the position-only against combined breakdown (right) for the five models. The single-agent and curriculum models cluster at the top; the self-play variants reach at most 16.7 %. The combined position-and-orientation gate is the discriminating objective.

The per-scene error plots for the two self-play variants, Figure 10, show the pattern behind these numbers. Both solve only a small number of the easiest scenes, and their orientation error remains above 1.4rad on almost every scene, far outside the 0.2rad tolerance. The self-play models never satisfy the combined gate on any rotation-requiring scene.




Figure 10: Per-scene error for TASP-dPose (left, 16.7 %) and ASP-dPose (right, 6.7 %). Both solve only the easiest scenes; orientation error remains far above tolerance on the rest.

5.3 RQ3: Skill Composition Under Self-Play

The failure of self-play on the combined gate is not a failure to learn the two skills; it is a failure to combine them. For the goal-solving agent, the training logs record the per-axis success rates alongside the combined-gate rate. These are shown in Figure 11 and summarised in Table 5. Under the potential-based reward, the agent reaches a position-only success rate of 54 % and an orientation-only rate of 45 %, each individually substantial. If the two skills were satisfied independently, their product would give a combined rate of about 24 to 25 %. The measured combined rate is about 10 %, roughly 2.5 times below that prediction. The two skills exist but are not satisfied at the same time, which is the precise sense in which self-play fails on this task.

Self-play variant	Alice validity	Position SR	Orientation SR	Combined SR
Original reward (2,048 envs)	89.3 %	1.3 %	7.1 %	0.07 %
PBRS, outcome-based	83.9 %	54.1 %	45.0 %	10.0 %
PBRS, time-based + penalty	90.5 %	54.1 %	46.3 %	8.5 %

Table 5: Training-time per-axis and combined success rates for the goal-solving agent under self-play. Under the improvement reward the agent learns almost nothing; under the potential-based reward it learns each axis to roughly 50 % but the combined gate remains near 10 %, well below the independence product of about 24 to 25 %. Alice’s goal-validity rate is high throughout.

Two further observations from Table 5 sharpen the diagnosis. Under the improvement reward, the agent learns essentially nothing: a position-only rate of 1.3 % and a combined rate of 0.07 %. The potential-based reward is what raises the per-axis skills to roughly 50 %, consistent with the GAE argument of Section 3.8. The goal-proposing agent, in turn, is not the bottleneck. Its goal-validity rate stays between 83 and 91 % across variants, so it reliably proposes valid goals. The obstacle is the goal-solving agent’s inability to satisfy the coupled criterion under the non-stationary goal distribution, not a breakdown in goal proposal.




Figure 11: Self-play training diagnostics. Left: the goal-proposing agent’s goal-validity rate, between 83 and 91 % across variants. Right: the goal-solving agent’s per-axis success rates against its combined-gate success rate. Each axis reaches roughly 50 % while the combined gate remains near 10 %.

This answers RQ3. Asymmetric self-play does enable the learning of the individual position and orientation skills, each to about 50 %. It does not, however, enable their combination: the combined success rate is roughly 2.5 times below what the independent skills would predict. The distinction matters for interpretation. It locates the difficulty in skill composition under a shifting goal distribution, rather than in the acquisition of the skills themselves.

5.4 Cross-Environment Confirmation

The ordering single-agent $\approx$ curriculum $\gg$ self-play holds in both simulators, which indicates it is a property of the methods rather than of Isaac Lab. Figure 12 shows the success rates in Isaac Lab and in `gym-pusht`, and Figure 13 shows the per-scene outcomes as a heatmap. In

both environments the single-agent and curriculum models dominate while the self-play model remains low.




Figure 12: Scene success rate in Isaac Lab and **gym-pusht** under the same 30-scene gate. The ordering of the methods is preserved across the two simulators.

One nuance qualifies the equivalence of single-agent and curriculum training. At the small **gym-pusht** batch of 960 transitions per update, the curriculum model reaches 50.0% while the plain single-agent model reaches only 10.0%, so the curriculum is a large improvement there. At the large Isaac Lab batch of 7,920 the two are equivalent, at 76.7% and 80.0%. The curriculum therefore substitutes for the shaped reward when the batch is small and gradient estimates are noisy. It adds nothing once the batch is large enough for stable estimates on both objectives at once. The value of the curriculum is thus regime-dependent, a point taken up in Chapter 6.




Figure 13: Per-scene outcomes across models and simulators. Colour encodes the number of successful attempts per scene out of 20. The combined position-and-orientation block is the region that only the single-agent and curriculum models on Isaac Lab solve consistently.

5.5 Qualitative Behaviour and the Off-Policy Baseline

Figure 14 shows keyframes of PPO-PBRS solving three representative scenes: a forward push, a lateral push, and a combined translate-and-rotate scene. The policy approaches the block along the object-relative offset, pushes along the decoded direction, and repeats until the combined gate is met. This behaviour matches the quantitative result of Section 5.1: position is controlled reliably, and the harder scenes are those requiring a large orientation change.




Figure 14: Keyframes of PPO-PBRS on three held-out scenes: a forward push (left), a lateral push (centre), and a combined translate-and-rotate scene (right).

The off-policy baseline provides an external point of calibration from a different algorithm family. Under the constrained project budget, the SAC with Hindsight Experience Replay baseline did not converge within the time available. It is therefore not reported as a competitive number. Its role and status are recorded so that the on-policy results are read against a concrete off-policy attempt rather than in isolation. Published planar-pushing success rates provide the external calibration in its place. Completing this baseline is listed among the directions for future work in Chapter 6.

6 Discussion and Conclusion

This chapter interprets the results, states the contributions against the three research questions, records the limitations of the study, and lists directions for future work.

6.1 Interpretation

OLD RESULTS — pending experiment set: old single-seed numbers and the old reward-vs-training framing; to be rewritten around the ASP-failure regime analysis.

The central result of this thesis is a priority ordering. For a contact-rich multi-objective planar-pushing task under a single-GPU budget, effort spent on principled reward design produces a larger and more reliable improvement than effort spent on structured training dynamics. The three research questions decompose this result into an efficiency claim, a comparison claim, and a mechanistic explanation.

The efficiency claim is direct. Potential-based shaping reduces the parallelism required for a given success rate by a factor of about four, from 2,048 to 528 environments. It also improves the combined-objective success rate by 35 percentage points over the earlier reward at the same gate. The reason was developed in Section 3.5. The potential-based reward is dense, correctly signed on every push, and bounded, whereas the improvement reward is negative for most useful pushes and has near-zero variance. The policy-invariance guarantee of Section 2.4 allows this dense signal to be added without changing the task. The result confirms empirically that the guarantee holds under stochastic contact physics.

The comparison claim is that neither a forced curriculum nor asymmetric self-play improves on the plain shaped single-agent model. The curriculum is statistically equivalent to the single-agent model at the large batch used in Isaac Lab, and self-play is far worse. This does not contradict the literature that motivates curricula and self-play; it locates a regime in which they do not pay. The cross-environment result of Section 5.4 makes the regime explicit. The curriculum substitutes for the shaped reward when the batch is small and gradient estimates are noisy, and adds nothing once the batch is large enough for stable estimates on both objectives. Reward design and curriculum design are therefore not independent levers of equal standing. At a sufficient batch size, the shaped reward already supplies what the curriculum would otherwise provide.

The mechanistic explanation concerns why self-play in particular fails. Section 5.3 showed that self-play learns the two constituent skills, each to roughly 50 %, but combines them at about 10 %, well below the independence prediction. Three factors act together to produce this outcome, and all are specific to this task. First, the physics couples translation and rotation, so the combined gate cannot be satisfied by treating the objectives separately. Second, the goal distribution is non-stationary, because the goal-proposing agent changes it at every iteration, which biases the value term of the GAE advantage as shown in Section 3.8. Third, the combined gate is strict, so partial competence on each axis does not accumulate into success. The potential-based reward supplies enough correctly-signed signal to learn each axis, but it cannot remove the value-term bias that prevents their composition. This is why shaping lifts self-play above the near-zero level it reaches under the improvement reward, yet leaves it far below single-agent training.

6.2 Relation to Prior Work

OLD RESULTS — pending experiment set: conceptual arguments are kept, but the embedded gym curriculum numbers (10.0% to 50.0%) are old single-seed values.

The findings of this thesis agree with the theory of the papers it builds on, while qualifying the empirical claims of the papers it tests. The distinction between the two is worth stating precisely.

On reward shaping, the result confirms [6] under conditions the original theorem did not examine. Ng, Harada, and Russell proved policy invariance for finite, discrete Markov Decision Processes with deterministic dynamics. The present task is continuous, high-dimensional, and driven by stochastic PhysX contact. The shaping potential is a bounded exponential, rather than the distance-to-goal heuristics of the original paper. The potential-based reward reaches the same success rate as the hand-tuned reward while requiring a quarter of the environments, and never exceeds it. This is the empirical signature the theorem predicts: shaping changes how quickly the optimal policy is found, not which policy is optimal. The episodic treatment of [19] licences the undiscounted shaping used here, and the stable value loss in Appendix B matches the bounded-potential construction. In this respect the thesis extends an established theoretical result into a contact-rich robotic regime rather than contesting it.

On asymmetric self-play the relationship is more nuanced, and the negative result must not be overstated. [8] demonstrated self-play in reversible or resettable gridworlds and continuous control tasks, where the goal is a single state to be reached. [9] scaled it to robotic manipulation with large distributed compute and a continuous action space. The present study differs from both on three axes that Sections 3.8 and 6.1 identify as decisive. The goal is multi-objective with a strict combined gate, not a single reachable state. The action is a discrete macro-action over a large MultiCategorical space, not a smooth continuous control. The compute budget is a single Graphical Processing Unit, not a distributed cluster. The finding is therefore not that self-play is ineffective. Rather, in this specific setting it does not compose the two learned skills, and it stays far below single-agent training under a budget realistic for one laboratory. This framing is consistent with [34], whose success with competitive self-play required months of training at very large scale. Through that lens, the modest self-play result here supports rather than contradicts the observation that self-play is demanding of scale.

The behaviour of the Alice Behavioral Cloning term deserves specific comment. It connects a design choice in [9] to an outcome observed here. Plappert and colleagues introduced Proximal Policy Optimization-style clipping on the behavioural-cloning loss (Equation 13), to prevent the imitation term from producing aggressive policy updates. In their continuous-action grasping domain, the demonstrated actions of the goal-proposing agent stay close to the goal-solving agent’s policy, so the clipping is rarely active and the imitation signal flows. In the discrete, contact-rich pushing domain of this thesis, the goal-proposing agent’s unguided pushes lie far from the goal-solving agent’s goal-directed policy. The probability ratio in Equation 13 is therefore extreme, the clip saturates, and the imitation gradient vanishes. The same mechanism that stabilises imitation in the original setting removes it here. This is not a flaw in the original formulation. It is a consequence of applying that formulation to a domain whose action distribution differs sharply from the one it was designed for. This also explains why the two self-play variants differ only through Alice’s reward, not through the cloning term.

On curriculum learning, the regime-dependence of Section 5.4 refines rather than rejects the

surveys of [7] and [25]. Those surveys motivate curricula primarily for settings where reward is sparse or exploration is hard. The present result is consistent with that motivation and makes it quantitative. The forced curriculum contributes substantially when the batch is small and gradient estimates are noisy: at the small **gym-pusht** batch it raises scene success from 10.0 % to 50.0 %. It contributes nothing once the batch is large enough for stable estimates on both objectives. A well-shaped dense reward and a curriculum are therefore partial substitutes, not independent and additive interventions. The point at which one ceases to add value over the other is governed by the gradient-variance regime. To the best of the author’s knowledge, this observation is not made explicit in the cited surveys, and it is one of the more transferable findings of the study.

Finally, the collapse of the symmetric time-based self-play variant to a fail-fast equilibrium (Appendix I) is best read through the game-dynamics analysis of [33]. In the symmetric reward, the goal-solving agent is penalised for the very solve time that rewards the goal-proposing agent. This defines a near-zero-sum game whose only stable point under the observed physics is immediate failure. The asymmetric reward of [8] incentivises the proposer without penalising the solver, and so avoids this degenerate equilibrium. The present study provides a concrete manipulation-domain instance of why that asymmetry is not incidental but necessary.

6.3 Contributions Revisited

OLD RESULTS — pending experiment set: stated against the old RQ1–RQ3 with old single-seed numbers; to be remapped to the three contributions of the ASP-failure framing.

The contributions stated in Chapter 1 can now be given quantitative content, one per research question.

For RQ1, the thesis formulates a potential-based reward for multi-objective planar pushing. It combines bounded exponential position and orientation potentials with a unified SE(2) distance. This reward reaches 80.0 % scene success at 528 environments, matching the improvement reward at 2,048 environments and improving the combined-objective rate by 35 percentage points.

For RQ2, the thesis provides a controlled comparison of four training regimes under an identical reward, budget, and evaluation protocol. The forced curriculum is statistically equivalent to the single-agent model, at 76.7 % against 80.0 %, a difference that is not significant. Both self-play variants are far worse, at 16.7 % and 6.7 %. The cross-environment comparison establishes that the value of the curriculum depends on the per-update batch size.

For RQ3, the thesis diagnoses the self-play failure by separating per-axis competence from combined success. Self-play learns each of position and orientation to roughly 50 %, but combines them at about 10 %, roughly 2.5 times below the independence product. The goal-proposing agent, meanwhile, remains a reliable source of valid goals. This locates the failure in skill composition under a non-stationary goal distribution, not in skill acquisition.

6.4 Limitations

OLD RESULTS — pending experiment set: states the single-seed / no-confidence-

interval limitation that the new experiment set (five seeds with confidence intervals) removes; must be updated.

Several limitations bound the scope of these conclusions. Validation uses a single seed per model, so the reported scene success rates carry no confidence intervals. Only the training-time success rate has variance data, across five chains for the single-agent model. The best-of-20 protocol raises the scene success rate above the trial success rate by 14 percentage points for the strongest model. Both numbers are reported, but the headline figure is the more favourable of the two by construction. The comparison between self-play and single-agent training differs on more than one axis at once. These include the number of agents, the goal distribution, the episode budget, the goal encoder, and the behavioural-cloning term. The comparison therefore isolates self-play as implemented, rather than the goal-proposal mechanism alone, and the axes of variation are enumerated in Appendix J. The sensitivity of the potential-based reward to its parameters k_p , k_r , and w was not systematically tested; the values used were chosen from pilot runs. Finally, the off-policy SAC with Hindsight Experience Replay baseline did not converge within the available budget, so the external calibration relies on published planar-pushing results rather than a converged in-house baseline.

6.5 Future Work

OLD RESULTS — pending experiment set: several items listed here (multi-seed confidence intervals, batch/scale sweep, ABC-clip instrumentation, disc isolation) are already performed by the experiment set; reconcile once its results exist.

Each limitation suggests a concrete continuation. Repeating the validation across multiple seeds would attach confidence intervals to the scene success rates. It would also place the 3.3-percentage-point single-agent-against-curriculum difference on a firmer statistical footing. A relaxed or partial-credit success gate, which rewards satisfying each axis rather than only the conjunction, would test whether the composition failure can be alleviated by softening the strict gate. The symmetric time-based self-play variant collapses to a fail-fast equilibrium (Appendix I); the zero-sum stability analysis of [33] suggests where a stabilised formulation might be sought. A systematic sweep of the per-update batch size would map the boundary of the regime in which the curriculum substitutes for the shaped reward, identified qualitatively in Section 5.4. Completing the SAC with Hindsight Experience Replay baseline to convergence would supply the off-policy calibration that the present study leaves open. Finally, transferring the shaped single-agent policy to a physical UR5e would test whether the efficiency gain survives the reality gap.

Beyond these direct continuations, the relation to prior work in Section 6.2 points to two deeper questions. The first is whether self-play can be made competitive by addressing the specific failure identified, rather than by adding scale. The value term of the Generalized Advantage Estimation advantage is biased by the non-stationary goal distribution (Section 3.8). A goal-proposing agent that changes the distribution more slowly might keep the value estimate accurate enough for the two skills to compose. One could constrain the rate of change of the proposed goal distribution, or train the goal-solving agent against a mixture of current and past proposers. This would test directly whether the composition failure is intrinsic to self-play or an

artefact of an over-aggressive proposer. The second question concerns the behavioural-cloning term. The clipping mechanism of [9] removes the imitation signal in the discrete action space used here. An unclipped or adaptively clipped cloning loss, or one defined on the decoded continuous push parameters rather than the discrete bins, might restore the demonstration signal the original formulation relies on. Both directions treat the negative self-play result not as a dead end, but as a localisation of where the method breaks in this domain.

6.6 Conclusion

OLD RESULTS — pending experiment set: old priority-ordering message and old numbers (4x, 35 percentage points); to be rewritten as the ASP-failure conclusion.

This thesis compared reward shaping against structured training dynamics on a single contact-rich planar-pushing task, under a common budget and a common success criterion. Potential-based reward shaping reduced the parallel-environment requirement by a factor of about four and improved combined-objective success by 35 percentage points. A forced curriculum added no value at a sufficient batch size, and asymmetric self-play performed far worse than single-agent training. The self-play failure was traced to an inability to combine two individually-learned skills under a non-stationary goal distribution, not to any failure to learn the skills themselves. The practical recommendation that follows is simple: invest first in a principled dense reward, and add structured training dynamics only in the small-batch regime where they demonstrably contribute.

Bibliography

- [1] M. Q. Mohammed, L. C. Kwek, S. C. Chua, A. Al-Dhaqm, S. Nahavandi, T. A. E. Eisa, M. F. Miskon, M. N. Al-Mhiqani, A. Ali, M. Abaker, and E. A. Alandoli, "Review of learning-based robotic manipulation in cluttered environments," *Sensors*, vol. 22, no. 20, p. 7938, 2022.
- [2] A. Zeng, S. Song, S. Welker, J. Lee, A. Rodriguez, and T. Funkhouser, "Learning synergies between pushing and grasping with self-supervised deep reinforcement learning," in *2018 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pp. 4238–4245, IEEE, 2018.
- [3] K. Mokhtar, C. J. M. Heemskerk, and H. Kasaei, "Self-supervised learning for joint pushing and grasping policies in highly cluttered environments," *ArXiv*, vol. abs/2203.02511, 2022.
- [4] H. Kasaei and M. M. Kasaei, "Harnessing the synergy between pushing, grasping, and throwing to enhance object manipulation in cluttered scenarios," 2024.
- [5] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms," *ArXiv*, vol. abs/1707.06347, 2017.
- [6] A. Y. Ng, D. Harada, and S. Russell, "Policy invariance under reward transformations: Theory and application to reward shaping," in *Proceedings of the Sixteenth International Conference on Machine Learning (ICML)*, pp. 278–287, 1999.
- [7] S. Narvekar, B. Peng, M. Leonetti, J. Sinapov, M. E. Taylor, and P. Stone, "Curriculum learning for reinforcement learning domains: A framework and survey," *arXiv preprint arXiv:2003.04960*, 2020.
- [8] S. Sukhbaatar, Z. Lin, I. Kostrikov, G. Synnaeve, A. Szlam, and R. Fergus, "Intrinsic motivation and automatic curricula via asymmetric self-play," in *International Conference on Learning Representations (ICLR)*, 2018.
- [9] O. OpenAI, M. Plappert, R. Sampedro, T. Xu, I. Akkaya, V. Kosaraju, P. Welinder, R. D'Sa, A. Petron, H. P. de Oliveira Pinto, A. Paino, H. Noh, L. Weng, Q. Yuan, C. Chu, and W. Zaremba, "Asymmetric self-play for automatic goal discovery in robotic manipulation," *ArXiv*, vol. abs/2101.04882, 2021.
- [10] M. T. Mason, "Mechanics and planning of manipulator pushing operations," *The International Journal of Robotics Research*, vol. 5, no. 3, pp. 53–71, 1986.
- [11] S. Goyal, A. Ruina, and J. Papadopoulos, "Planar sliding with dry friction part 1. limit surface and moment function," *Wear*, vol. 143, no. 2, pp. 307–330, 1991.
- [12] M. Mittal, C. Yu, Q. Yu, J. Liu, N. Rudin, D. Hoeller, J. L. Yuan, R. Singh, Y. Guo, H. Mazhar, A. Mandlekar, B. Babich, G. State, M. Hutter, and A. Garg, "Orbit: A unified simulation framework for interactive robot learning environments," *IEEE Robotics and Automation Letters*, vol. 8, no. 6, pp. 3740–3747, 2023.
- [13] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*. The MIT Press, second ed., 1998.

-
- [14] M. Yang, Y. Lin, A. Church, J. Lloyd, D. Zhang, D. A. W. Barton, and N. F. Lepora, “Sim-to-real model-based and model-free deep reinforcement learning for tactile pushing,” *IEEE Robotics and Automation Letters*, vol. 8, pp. 5480–5487, 2023.
 - [15] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. A. Riedmiller, A. K. Fidjeland, G. Ostrovski, S. Petersen, C. Beattie, A. Sadik, I. Antonoglou, H. King, D. Kumaran, D. Wierstra, S. Legg, and D. Hassabis, “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, pp. 529–533, 2015.
 - [16] T. Haarnoja, A. Zhou, P. Abbeel, and S. Levine, “Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor,” *ArXiv*, vol. abs/1801.01290, 2018.
 - [17] J. Schulman, S. Levine, P. Abbeel, M. I. Jordan, and P. Moritz, “Trust region policy optimization,” *ArXiv*, vol. abs/1502.05477, 2015.
 - [18] M. Andrychowicz, F. Wolski, A. Ray, J. Schneider, R. Fong, P. Welinder, B. McGrew, J. Tobin, O. Abbeel, and W. Zaremba, “Hindsight experience replay,” *Advances in Neural Information Processing Systems*, vol. 30, 2017.
 - [19] M. Grześ, “Reward shaping in episodic reinforcement learning,” in *Proceedings of the 16th Conference on Autonomous Agents and MultiAgent Systems (AAMAS)*, pp. 565–573, 2017.
 - [20] S. Devlin and D. Kudenko, “Dynamic potential-based reward shaping,” in *Proceedings of the 11th International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, pp. 433–440, 2012.
 - [21] A. Harutyunyan, S. Devlin, P. Vrancx, and A. Nowé, “Expressing arbitrary reward functions as potential-based advice,” in *Proceedings of the Twenty-Ninth AAAI Conference on Artificial Intelligence*, pp. 2652–2658, 2015.
 - [22] M. Grześ and D. Kudenko, “Theoretical and empirical analysis of reward shaping in reinforcement learning,” in *2009 International Conference on Machine Learning and Applications (ICMLA)*, pp. 337–344, IEEE, 2009.
 - [23] C. Florensa, D. Held, M. Wulfmeier, M. Zhang, and P. Abbeel, “Reverse curriculum generation for reinforcement learning,” in *Conference on Robot Learning (CoRL)*, pp. 482–495, PMLR, 2017.
 - [24] S. Luo, H. Kasaei, and L. Schomaker, “Accelerating reinforcement learning for reaching using continuous curriculum learning,” *2020 International Joint Conference on Neural Networks (IJCNN)*, pp. 1–8, 2020.
 - [25] R. Portelas, C. Colas, L. Weng, K. Hofmann, and P.-Y. Oudeyer, “Automatic curriculum learning for deep rl: A short survey,” *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 4819–4825, 2020.
 - [26] F. Torabi, G. Warnell, and P. Stone, “Behavioral cloning from observation,” in *Proceedings of the Twenty-Seventh International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 4950–4957, 2018.

- [27] P. Florence, C. Lynch, A. Zeng, O. A. Ramirez, A. Wahid, L. Downs, A. Wong, J. Lee, I. Mordatch, and J. Tompson, “Implicit behavioral cloning,” in *Conference on Robot Learning (CoRL)*, pp. 158–168, PMLR, 2022.
- [28] T. Hester, M. Vecerík, O. Pietquin, M. Lanctot, T. Schaul, B. Piot, D. Horgan, J. Quan, A. Sendonaris, I. Osband, G. Dulac-Arnold, J. P. Agapiou, J. Z. Leibo, and A. Gruslys, “Deep q-learning from demonstrations,” in *AAAI Conference on Artificial Intelligence*, 2017.
- [29] T. Silver, K. Allen, J. Tenenbaum, and L. Kaelbling, “Residual policy learning,” *arXiv preprint arXiv:1812.06298*, 2018.
- [30] M. Dennis, N. Jaques, E. Vinitzky, A. Bayen, S. Russell, A. Critch, and S. Levine, “Emergent complexity and zero-shot transfer via unsupervised environment design,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 13049–13061, 2020.
- [31] A. Campero, R. Raileanu, H. Küttler, J. B. Tenenbaum, T. Rocktäschel, and E. Grefenstette, “Learning with amigo: Adversarially motivated intrinsic goals,” in *International Conference on Learning Representations (ICLR)*, 2021.
- [32] I. Durugkar, M. Tec, S. Niekum, and P. Stone, “Adversarial intrinsic motivation for reinforcement learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, pp. 8622–8636, 2021.
- [33] A. Letcher, D. Balduzzi, S. Racanière, J. Martens, J. Foerster, K. Tuyls, and T. Graepel, “Differentiable game mechanics,” *Journal of Machine Learning Research*, vol. 20, no. 84, pp. 1–40, 2019.
- [34] C. Berner, G. Brockman, B. Chan, V. Cheung, P. Debiak, C. Dennison, D. Farhi, Q. Fischer, S. Hashme, C. Hesse, *et al.*, “Dota 2 with large scale deep reinforcement learning,” *arXiv preprint arXiv:1912.06680*, 2019.
- [35] S. Sukhbaatar, E. Denton, A. Szlam, and R. Fergus, “Learning goal embeddings via self-play for hierarchical reinforcement learning,” in *arXiv preprint arXiv:1811.09083*, 2018.
- [36] N. Tishby and N. Zaslavsky, “Deep learning and the information bottleneck principle,” in *2015 IEEE Information Theory Workshop (ITW)*, pp. 1–5, IEEE, 2015.
- [37] A. Goyal, R. Islam, D. Strouse, Z. Ahmed, H. Larochelle, M. Botvinick, S. Levine, and Y. Bengio, “Infobot: Transfer and exploration via the information bottleneck,” in *International Conference on Learning Representations (ICLR)*, 2019.
- [38] O. Nachum, S. S. Gu, H. Lee, and S. Levine, “Data-efficient hierarchical reinforcement learning,” in *Neural Information Processing Systems*, 2018.
- [39] A. S. Vezhnevets, S. Osindero, T. Schaul, N. Heess, M. Jaderberg, D. Silver, and K. Kavukcuoglu, “Feudal networks for hierarchical reinforcement learning,” in *International Conference on Machine Learning (ICML)*, pp. 3540–3549, PMLR, 2017.
- [40] P.-L. Bacon, J. Harb, and D. Precup, “The option-critic architecture,” in *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence*, pp. 1726–1734, 2017.

- [41] M. Hutsebaut-Buysse, K. Mets, and S. Latré, “Hierarchical reinforcement learning: A survey and open research challenges,” *Machine Learning and Knowledge Extraction*, vol. 4, no. 1, pp. 172–221, 2022.
- [42] A. V. Nair, V. Pong, M. Dalal, S. Bahl, S. Lin, and S. Levine, “Visual reinforcement learning with imagined goals,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31, 2018.
- [43] K. M. Lynch and M. T. Mason, “Stable pushing: Mechanics, controllability, and planning,” *The International Journal of Robotics Research*, vol. 15, no. 6, pp. 533–556, 1996.
- [44] R. D. Howe and M. R. Cutkosky, “Practical force-motion models for sliding manipulation,” *The International Journal of Robotics Research*, vol. 15, no. 6, pp. 557–572, 1996.
- [45] J. Stüber, C. Zito, and R. Stolkin, “Let’s push things forward: A survey on robot pushing,” *Frontiers in Robotics and AI*, vol. 7, p. 8, 2020.
- [46] K.-T. Yu, M. Bauza, N. Fazeli, and A. Rodriguez, “More than a million ways to be pushed: A high-fidelity experimental dataset of planar pushing,” *2016 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pp. 30–37, 2016.
- [47] F. C. Park, J. E. Bobrow, and S. R. Ploen, “A lie group formulation of robot dynamics,” *The International Journal of Robotics Research*, vol. 14, no. 6, pp. 609–618, 1995.
- [48] J. Urain, N. Funk, J. Peters, and G. Chalvatzaki, “Se(3)-diffusionfields: Learning smooth cost functions for joint grasp and motion optimization through diffusion,” in *2023 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 5923–5930, IEEE, 2023.
- [49] E. van der Pol, D. E. Worrall, H. van Hoof, F. A. Oliehoek, and M. Welling, “Mdp homomorphic networks: Group symmetries in reinforcement learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 4199–4210, 2020.
- [50] H. Huang, D. Wang, R. Walters, and R. Platt, “Equivariant transporter network,” in *Robotics: Science and Systems (RSS)*, 2022.
- [51] H. Nguyen, A. Baisero, D. Klee, D. Wang, R. Platt, and C. Amato, “Equivariant reinforcement learning under partial observability,” in *Conference on Robot Learning (CoRL)*, 2024.
- [52] V. Makoviyshuk, L. Wawrzyniak, Y. Guo, M. Lu, K. Storey, M. Macklin, D. Hoeller, N. Rudin, A. Allshire, A. Handa, and G. State, “Isaac gym: High performance gpu-based physics simulation for robot learning,” in *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2021.
- [53] P. Henderson, R. Islam, P. Bachman, J. Pineau, D. Precup, and D. Meger, “Deep reinforcement learning that matters,” in *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence*, pp. 3207–3214, 2018.
- [54] R. Agarwal, M. Schwarzer, P. S. Castro, A. C. Courville, and M. G. Bellemare, “Deep reinforcement learning at the edge of the statistical precipice,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, pp. 29304–29320, 2021.

Appendices

A Push Mechanics: Limit Surface and Characteristic Length

The limit-surface normality condition of Equation 14 states that the instantaneous motion twist of a quasi-statically pushed object is normal to its friction limit surface at the point corresponding to the applied wrench. Its practical consequence for the task concerns any off-centre push, one whose line of action does not pass through the object’s centre of friction. Such a push produces a wrench with a non-zero moment component, and therefore a twist that contains an angular velocity. Almost every push consequently changes both position and orientation. The characteristic length L used in the SE(2) distance of Equation 20 is the radius of gyration of the object, a physical constant; for the T-block $L = 0.07\text{ m}$. It is not a tuning parameter, and it sets the weight of an orientation error relative to a position error in a manner consistent with the object’s mass distribution. For the rotationally symmetric disc the characteristic length is set to zero, which collapses the SE(2) distance to a pure position distance and encodes the absence of an orientation objective. The mechanical contrast between the two objects, the coupled twist of the T-block against the pure translation of the disc, is discussed in Section 2.8.




Figure 15: An off-centre push produces coupled translation and rotation of the T-block. The motion twist is normal to the friction limit surface.

B Training Curves

Figure 16 shows the training-time success rate and value loss for the single-agent and curriculum models. The single-agent model improves steadily over roughly 3,000 iterations, the curriculum model follows a similar trajectory, and the value loss remains stable throughout, which reflects the bounded range of the exponential potentials.




Figure 16: Training success rate (left) and value loss (right) for the potential-based models.

C PBRS Parameter Sensitivity

The potential-based reward uses $k_p = 30$, $k_r = 5$, and $w_{\text{pos}} = w_{\text{rot}} = 10$, chosen from pilot runs. A systematic grid over these parameters was not performed within the available budget. The robustness of the reward to its parameters is therefore not established here. It is recorded as a limitation in Section 6.4 and as a direction for future work in Section 6.5. A grid over $k_p \in \{10, 20, 30, 50\}$ and $w \in \{5, 10, 20\}$, evaluated at the 30-scene gate, would provide the missing sensitivity analysis.

D Validation Scene Descriptions

The 30-scene held-out suite is defined in `validation_configs.py` and comprises three groups of ten. Scenes 1 to 10 are rotation-dominant, with orientation changes up to π radians and mixed translations. Scenes 11 to 20 are position-only, with the goal orientation equal to the start orientation. Scenes 21 to 30 combine a position change with an orientation change. Within each group, difficulty is graded easy, medium, and hard by the start-to-goal distance and the magnitude of the required rotation. All coordinates are expressed in the environment-local frame in metres, and orientations in radians.

E cuRobo IK Pipeline and Workspace

Each decoded push primitive is executed as a five-phase Cartesian trajectory, approach, descend, push, retract, and return, spanning 72 physics substeps. Every waypoint is mapped to arm joint targets by the cuRobo 0.7.5 batched Inverse Kinematics solver, with the gripper held closed throughout. The reachable workspace is bounded to the region over the table where the solver returns elbow-up solutions. Pushes whose start or end points fall outside this region are clamped to the boundary before execution, so the credit assigned to an action matches the motion actually performed.

F Observation Space Specifications

The single-agent models use a 28-dimensional observation: end-effector pose (6), object pose and velocity (14), goal pose (6), and the position and orientation errors (2). The goal-solving agent in the self-play models uses the same 28-dimensional layout, with the goal encoded through the GoalEncoder into an eight-dimensional latent in the difference variant. The goal-proposing agent uses a 20-dimensional observation that omits the goal, since it sets the goal rather than solving it.

G Per-Test Error Bars

OLD RESULTS — pending experiment set: single-seed per-scene bars and the “statistically equivalent” claim; to be regenerated with the seed-CI data.

Figure 17 shows the per-scene trial-level success rate with 95 % binomial confidence intervals for the single-agent and curriculum models head to head, and for all models together. The confidence intervals of the single-agent and curriculum models overlap on most scenes, consistent with the finding of Section 5.2 that the two are statistically equivalent.




Figure 17: Per-scene trial-level success rate with 95 % binomial confidence intervals.

H Curriculum Trigger Correction

An earlier implementation of the forced curriculum gated the transition to the orientation phase on the per-push mean position error falling below 0.08 m for 50 consecutive iterations. This condition never activated. The per-push mean is dominated by early-episode pushes far from a freshly randomised goal, and has a structural floor above 0.08 m. Even the best model’s per-push mean position error does not fall below this value, and a single noisy iteration reset the 50-iteration streak. The orientation weight therefore remained at zero for the entire run, and the model reduced to a position-only learner. The corrected trigger, used for the PPO-Curriculum results in this thesis, gates instead on the episodic position-only success rate reaching 0.5 over a 20-iteration window, which is a “position is solved” signal rather than a per-push error mean, and which activates as intended.




Figure 18: The original curriculum trigger never activated: the per-push mean position error remained above the 0.08 m threshold throughout training.

I Symmetric Self-Play Fail-Fast Equilibrium

OLD RESULTS — pending experiment set: the 0% collapse is an old single-seed observation; the Bob-penalty phase of the experiment set will quantify it with confidence intervals.

A symmetric variant of time-based self-play, in which the goal-solving agent is penalised for its own solve time in addition to the goal-proposing agent being rewarded for it, collapses to a 0% success rate. The penalty applies on all phase ends, including catastrophic ones, so the agent obtains a higher return by failing on the first push than by attempting a longer solution and failing later. Because failing is easier than succeeding under contact-rich physics, the policy converges to immediate failure, and the goal-proposing agent’s reward collapses to zero at the same time. The lesson is that an asymmetric reward structure is necessary in this domain: the goal-proposing agent must be incentivised, and the goal-solving agent must not be penalised for time. This is a reward-design outcome specific to the setting, not a refutation of the symmetric formulation in general.

J Confounded Ablation Disaggregation

OLD RESULTS — pending experiment set: the "10-percentage-point" advantage is an old number, and the new single-axis isolations (disc, ABC on/off, encoder on/off) change this account.

The single-agent and self-play models differ on at least seven axes at once: the number of agents (one against two), the goal distribution (fixed random against adversarial and shifting), the goal encoder (absent against present), the behavioural-cloning term (absent against present), the historical opponent pool (absent against present), the episode budget (15 pushes for one agent against 5 plus 10 split between two), and the number of policy updates per iteration (one against two). The comparison in this thesis therefore isolates self-play as a package rather than any single mechanism within it. The one clean single-axis comparison available is between

the two self-play variants, which differ only in the goal-proposing agent’s reward and show a 10-percentage-point advantage for the time-based formulation.

K Training Budget Comparison

OLD RESULTS — pending experiment set: old single-seed 80.0% figures and the “five principal models” framing.

The single-agent and curriculum models and both self-play models were trained at 528 environments for approximately 3,000 iterations. The PPO-Baseline was trained at 2,048 environments, a deliberately larger batch, so that its lower combined-objective performance could not be attributed to insufficient parallelism. Both the baseline and the potential-based single-agent model reach 80.0 % scene success, but the potential-based model does so with a quarter of the parallel environments, which is the environment-efficiency result of Section 5.1.

L Secondary Variants

OLD RESULTS — pending experiment set: treats disc and Bob-penalty as “secondary” and cites old results; these are now first-class subjects in the experiment set.

Beyond the five principal models, several secondary variants were trained and evaluated. These are self-play on a rotationally symmetric disc, for which the $SE(2)$ distance reduces to pure position distance; the original, non-simplified potential formulation at the position-only gate; and the non-HPC `gym-pusht` runs at reduced core counts. All are consistent with the principal findings: self-play remains far below single-agent training across every one of them. They are omitted from the main text for concision.